

COGNITIVE DIGNITY: A PROPOSAL FOR AN EMERGING HUMAN RIGHT

Protecting the Human Process of Sense-Making and Knowledge Formation against Technological Substitution

Alejandro Sinopoli

Independent Researcher · Rosario, Argentina

epistepolis@gmail.com

ORCID: 0009-0005-7790-8476

Abstract: This paper presents Cognitive Dignity as a doctrinal proposal for an emerging human right aimed at protecting the person as the active subject of their own process of sense-making and knowledge formation. Its object is not merely the content of thought, mental data, neural integrity, or the final decision, but the process itself through which a person thinks, understands, revises, integrates, and attributes meaning. The distinctive problem arises when technological or institutional architectures materially supplant, foreclose, or degrade that process, even in the absence of brain reading, psychological manipulation, deception, or coercion. The paper compares the proposal, source by source, with freedom of thought, cognitive liberty and neurorights, the right to mental integrity, human dignity in the face of artificial intelligence, epistemic rights, mental self-determination, and epistemic agency, and tests it against a concrete judicial case (Colombian Constitutional Court, Judgment T-323 of 2024). No claim is made to terminological originality or universal priority; within the audited corpus (Annex A), no single earlier source has been identified that brings together the same complete legal-conceptual formula. The analysis concludes that there is a residual gap and a defensible material increment centered on non-manipulative cognitive substitution, although the need to recognize a free-standing autonomous human right distinct from the evolution of dignity, freedom of thought, autonomy, and cognitive liberty has not been demonstrated.

Keywords: cognitive dignity; artificial intelligence; human rights; freedom of thought; cognitive liberty; neurorights; mental integrity; mental self-determination; epistemic agency; cognitive substitution; Habeas Mens.

Scope note. This paper advances a doctrinal proposal and a hypothesis of specific originality within the audited corpus. It does not claim that Cognitive Dignity is today an autonomous human right recognized by international law, nor that absolute universal priority has been demonstrated.

Suggested citation: Sinopoli, A. (2026). Cognitive Dignity: A Proposal for an Emerging Human Right. Protecting the Human Process of Sense-Making and Knowledge Formation against Technological Substitution (Working Paper, English Version 1.7). SSRN.

1. Introduction

The expansion of artificial intelligence systems capable of producing text, diagnoses, recommendations, evaluations, and outputs that perform or simulate reasoning functions has shifted a significant part of legal debate from access to information toward the distribution of cognitive functions between persons and systems. Existing frameworks provide robust responses to multiple risks: privacy, data processing, discrimination, manipulation, autonomy, mental integrity, freedom of thought, due process, and human oversight. Yet a different question remains: what happens when no one reads a person's mind, no one deliberately alters their beliefs, and no one coerces them, but the process through which they would have understood, worked through, or produced meaning is materially supplanted by an external architecture?

Cognitive Dignity arises as a response to that problem. Its central thesis is not that all technological assistance is harmful, nor that legitimate knowledge must be produced without tools. The claim is narrower: there are contexts in which a person's status as a knowing subject has normative significance in its own right and may be degraded even when classical rights concerning mental contents, privacy, or the final decision have not clearly been violated.

The concept developed through two documented moves. On 21 October 2025, the formulation “preserve cognitive dignity in the face of the automation of knowledge” appears, together with an epistemic-responsibility rule according to which “no knowledge is legitimate if it degrades the dignity of the knower.” On 5 December 2025, the idea is expressly framed in legal terms as an “emerging human right” aimed at protecting the inner or symbolic process through which a person thinks, feels, understands, and makes sense. This article takes that formulation as its point of departure and refines it legally without retrospectively attributing to 2025 all later developments. (Sinopoli, 2025c, 2025d).

The argument is organized around four claims. First, Cognitive Dignity has a specific object: the human process of sense-making and knowledge formation. Second, that object overlaps with prior rights but is not fully reducible to privacy, mental integrity, cognitive liberty, or decisional autonomy. Third, non-manipulative cognitive substitution identifies the clearest differential case. Fourth, the existence of that residue is not yet sufficient to demonstrate that an autonomous human right must be created: Cognitive Dignity may also operate as a doctrinal specification of existing rights. Section 7.2 tests this thesis against a concrete judicial case, and Annex A describes the search protocol supporting the claims concerning antecedents.

2. Documentary Genealogy and Temporal Delimitation

The history of the concept must be separated from the question of its universal originality. Registration certificates establish that a particular formulation existed on a given date; they do not by themselves prove that no one had formulated it earlier. This distinction is essential to avoid turning an author's chronology into an argument for scientific priority.

Date	Documentary milestone	Relevant content
17/09/2025	Epistemología Simbólica Comparativa [Comparative Symbolic Epistemology]	Verifiable beginning of the documentary genealogy. It does not yet contain the legal formula of Cognitive Dignity. (Sinopoli, 2025a)
15/10/2025	ESC–TESA–OCD	Documented consolidation of the prior architecture. (Sinopoli, 2025b)
21/10/2025	ESC TESA OCD v2.3	“Cognitive dignity in the face of the automation of knowledge” and the dignity of the knower as a criterion of epistemic responsibility. (Sinopoli, 2025c)
05/12/2025	EL TRATADO – EPISTEMÍA [THE TREATISE – EPISTEMÍA]	Explicit formulation as an emerging human right aimed at protecting the inner/symbolic process. (Sinopoli, 2025d)
20/12/2025	Consolidated Treatise	Further development and differentiation from privacy, cognitive liberty, decisional autonomy, and neural data. (Sinopoli, 2025e)
06/05/2026	Habeas Mens	Later legal-operational development concerning attributability and integrity of the human decision-making process; it is not synonymous with Cognitive Dignity. (Sinopoli, 2026a)

Accordingly, this paper uses two cutoff dates: 21 October 2025 for the epistemic-protective core and 5 December 2025 for the explicit formulation as an emerging human right. Earlier milestones establish intellectual continuity, but they must not be used to backdate a formulation that had not yet been expressed. (Sinopoli, 2025a, 2025b, 2025c, 2025d).

3. The Concept of Cognitive Dignity

3.1. Historical Formulation

“Cognitive Dignity is the emerging human right to the protection of the inner symbolic process.” (Sinopoli, 2025d, p. 27).

The December 2025 formulation specified that the right was not exhausted by mental privacy, freedom of thought, or the protection of neural data. Its object was “the very process by which we think, feel, and understand.” The text identified as risks substitution by prefabricated answers, acceleration beyond the person's own rhythms, invasion by exogenous content, premature closure, and the extraction of symbols without consent. (Sinopoli, 2025d, pp. 27-30).

3.2. Legally Refined Formulation

Cognitive Dignity is the proposal for an emerging human right aimed at preserving the person's status as the active subject of their own process of sense-making and knowledge formation against human, institutional, or technological architectures that materially supplant, foreclose, or degrade that process, including cases of non-manipulative cognitive substitution.

In this article, the expression “emerging human right” is not used as a technical category of positive law or as a claim of current international recognition. It is used in a doctrinal-programmatic sense: a claim linked to social or technological transformations, aimed at formulating a new right or renewing the scope of an existing one. This usage is consistent with the tradition of the Universal Declaration of Emerging Human Rights, whose final text was approved in Monterrey in 2007 and which the Institut de Drets Humans de Catalunya characterizes as a programmatic civil-society instrument oriented toward the crystallization of human rights in the new millennium (Institut de Drets Humans de Catalunya, 2009). Accordingly, Cognitive Dignity is presented here as a doctrinal proposal for an emerging right; no already-consolidated judicial, legislative, or institutional reception is presumed. Any such reception would constitute subsequent evidence of consolidation, not a semantic prerequisite of the proposal.

Nor does the use of the term “emerging” by itself transform a claim into a human right or replace the examination of necessity, non-redundancy, and precision. For that reason, Section 8 subjects the proposal to quality-control criteria inspired by Alston (1984) and Nickel (2007), while leaving open the possibility that protection may ultimately be accommodated, through evolutionary interpretation, within already recognized rights.

The legal refinement introduces limits that the early formulations did not make sufficiently explicit. Not every external intervention is unlawful; teaching, advising, assisting, translating, calculating, summarizing, or enhancing capacities may be entirely compatible with human dignity. The legal threshold must attend to the materiality of substitution, context, the existence of consent, available safeguards, reversibility, the possibility of challenge, and the proportionality of the architecture employed.

The reference to “human or institutional” architectures does not extend the concept to every interpersonal interaction. A teacher who offers an answer, a colleague who gives advice, or an expert who provides an opinion does not, by that fact alone, violate Cognitive Dignity. The reference concerns architectures: organizational, procedural, or normative designs that stably assign to an external system—technical or human—the cognitive function that the same environment continues to attribute to the person. The paradigmatic case, and the one motivating this paper, is technological; the inclusion of institutional architectures prevents

the same displacement from falling outside the concept merely because it is carried out through human rules, deadlines, or protocols.

3.3. Structural Elements

Dimension	Content of the proposal
Protected subject	The human person as an active subject of sense-making and knowledge formation.
Immediate object	The human process of thinking, understanding, revising, integrating, and making sense; not merely mental content or the final result.
Relevant properties	Temporality, openness, one's own elaboration, capacity for revision, and responsible appropriation.
Typical threats	Substitution, foreclosure, compression, or material degradation of the process; manipulation is a possible modality, but not a necessary one.
Status	Doctrinal proposal for an emerging human right.
Relationship to other rights	It overlaps with dignity, freedom of thought, autonomy, cognitive liberty, mental integrity, and sector-specific rights, but posits a differentiated processual object.
Differential case	Non-manipulative cognitive substitution: the human process is displaced without any need to read, alter, or manipulate the mind.

4. Method of Comparison: Avoiding Retrospective Collage

The originality of a conceptual formula cannot be refuted by assembling scattered fragments from multiple sources and treating them as though a single source had formulated the complete concept. The fact that one author addressed dignity, another epistemic agency, another automation, and another legitimacy shows that relevant intellectual materials were available; it does not by itself show that Cognitive Dignity had already been formulated.

Accordingly, each candidate is assessed individually by reference to: protected subject, protected object, threat, mechanism of harm, normative status, correlative duty, relationship to existing rights, and scope. Partial overlaps are expressly acknowledged, but they are not converted into conceptual identity by aggregation.

This method distinguishes four categories: conceptual anticipation; near anticipation; partial antecedent; and terminological antecedent or remote analogy. The decisive question is not

whether a source uses similar words, but whether it substantially protects the same object against the same type of threat with an equivalent normative function.

5. Position Relative to Prior Developments

5.1. Terminological Antecedents

The expression “cognitive dignity” is not original. Fornet-Betancourt (2007) used a related family of terms in intercultural philosophy to denounce the epistemic subordination of cultural traditions; Sorum (2017) used “cognitive dignity” in the context of medical diagnostic reasoning. These antecedents are sufficient to rule out any claim of priority over the words themselves, but their objects differ from the automation of the human process of knowing.

5.2. Cognitive Liberty and Neurorights

Cognitive liberty constitutes an earlier and closely related legal family. Sententia (2004) and Bublitz (2013) developed mental self-determination in relation to technologies capable of affecting cognition. Ienca and Andorno (2017) proposed rights or neurorights associated with cognitive liberty, mental privacy, mental integrity, and psychological continuity. These constructions occupy nearby territory, but they focus predominantly on access, alteration, manipulation, the use of neurotechnologies, or control over mental states. Cognitive Dignity shifts the focus toward the human process of sense-making and knowledge formation, including the case in which that process is simply supplanted externally without intervention in mental content.

Within that family, the most precise and recent formulation of the right to mental integrity is Cassinadri's (2025), which reconstructs it as a multidimensional right and expressly extends it to the non-neural extended mind. The test proposed in that account substantially requires a non-consensual interference that bypasses the person's rational capacities and produces mental harm. Precisely for that reason, it sharply delineates the residue of Cognitive Dignity: the cognitive substitution at issue here is often consented to or formally accepted, does not bypass reason through covert mechanisms, and, when it produces deterioration, does so gradually or longitudinally rather than through an identifiable episodic harm. The right to mental integrity thus reconstructed does not reach that case, which makes it not a rival but a complementary regime.

The neurorights family also includes positive-law and regional developments that a paper written from Latin America cannot omit. Through Law 21,383 (2021), Chile incorporated into Article 19(1) of its Constitution the requirement that scientific and technological development respect physical and psychological integrity and specifically safeguard brain activity and the information derived from it; the Supreme Court applied that clause in *Girardi Lavín v. Emotiv Inc.* (Case No. 105.065-2023, 9 August 2023) and ordered the deletion of brain data collected by a commercial device. In 2023, the Inter-American Juridical Committee adopted the Inter-American Declaration of Principles regarding Neuroscience, Neurotechnologies, and Human Rights, and the United Nations Human Rights Council Advisory Committee examined the effects of neurotechnology on human rights (A/HRC/57/61, 2024). These instruments confirm the vitality of this neighboring field, but their stated object is neurotechnology: brain activity and neural data. They do not address, as

an object in its own right, the non-neural substitution of the process of understanding by generative artificial-intelligence systems.

5.3. Freedom of Thought: The Strongest Legal Competitor

The most serious legal competitor is not a neuroright but freedom of thought under Article 18 of the International Covenant on Civil and Political Rights, whose internal dimension (forum internum) is absolute. After decades of limited development, that dimension has been intensely reinterpreted. McCarthy-Jones (2019) proposed reading it for the twenty-first century as protection of the autonomous mind against technologies capable of inferring and altering thoughts; Alegre (2022) defended it as an absolute right against commercial and political manipulation and inference; and the United Nations Special Rapporteur on freedom of religion or belief devoted his 2021 report to freedom of thought (A/76/380). The report identifies four possible attributes: not being compelled to reveal one's thoughts, not being penalized for them, not being subjected to impermissible alterations, and States' promotion of an enabling environment for free thought (A/76/380, para. 25); paragraphs 40-44 specifically develop that fourth attribute and its limits.

This is where Cognitive Dignity faces the greatest risk of being subsumed. If freedom of thought protects the process of forming thought rather than only its content, and if it includes a duty to preserve an enabling environment, an evolutionary interpretation could cover a substantial part of cognitive substitution. Two differences nevertheless remain. First, freedom-of-thought doctrine has been built around impermissible alteration—manipulation, coercion, inference—that is, interventions acting upon thought; non-manipulative substitution does not alter the person's thought but instead renders it unnecessary in the context. Second, the absolute character of the forum internum fits poorly with a case that requires contextual balancing (consent, function, severity, reversibility): forcing such cases into an absolute right would either trivialize its threshold or leave most cases outside it. Cognitive Dignity may therefore be read as a balanceable specification adjacent to the forum internum, rather than as its duplicate. This conclusion is provisional and depends on the interpretive evolution of Article 18.

5.4. Human Dignity, Artificial Intelligence, and Cognitive Autonomy

Boon (2021), in an academic-outreach text associated with an exhibition at the University of Twente, raised a direct concern about human dignity in the face of the possibility that artificial intelligence could replace intellectual capacities. Teo (2023) examined the usefulness of human dignity in responding to challenges posed by artificial intelligence, including cognitive autonomy. These are important antecedents that prevent an exaggerated narrative of novelty: the relationship among AI, dignity, and intellectual autonomy clearly predates this proposal. However, none of those formulations, as identified in the audited corpus, makes the human process of sense-making and knowledge formation the immediate object of an emerging right.

5.5. Epistemic Rights and the Dignity of the Knower

Bhakuni (2023) is relevant because, from a human-rights perspective, Bhakuni proposes a dignity-based approach to repairing epistemic wrongs and injustices, treating certain wrongs to the person as knower as injuries to dignity. This line confirms that the dignity of the epistemic subject does not originate with Cognitive Dignity. Its problem, however, lies in

epistemic injustice and recognition of the subject as a producer and user of knowledge, not in the automation or technological substitution of the process of knowing.

5.6. Faraoni (2023): The Closest Legal Antecedent

Among individual doctrinal proposals, the most demanding legal antecedent identified is Faraoni (2023); the strongest normative competitor, however, remains freedom of thought (§5.3). Faraoni's proposal for a right to mental self-determination against persuasive technologies and computational manipulation distinguishes between possessing or expressing a thought and the process of creating it; it also protects control over the decision-making process. This prevents any broad claim to originality of the form “there is a right for AI not to interfere with how we think.”

The difference lies in the core mechanism. Faraoni starts from persuasion, hypernudging, exploitation of vulnerabilities, and manipulation of decision formation. Cognitive Dignity includes that territory but adds a different case: substitution or foreclosure of the human process even when the system does not seek to modify mental content or induce a decision. An artificial response may be correct, transparent, and non-manipulative and yet displace the process of understanding that an institutional or educational architecture presupposes to be human. That difference does not by itself establish an autonomous right, but it does preserve a non-identical conceptual residue.

5.7. Epistemic Agency and Automation

Malone, Afroogh, D’Cruz, and Varshney (2025) connect automation, human dignity, epistemic harm, and diminished agency of expert workers. Coeckelbergh (2025) analyzes how AI may diminish epistemic agency in belief formation and revision. Both are close antecedents to the mechanism of cognitive displacement. Their focus, however, lies in epistemic agency, trust, expert work, or belief formation, rather than in a general right to the integrity of the human process of sense-making.

5.8. UNESCO and Institutional Frameworks

UNESCO's Guidance for Generative AI in Education and Research (2023) adopts a human-centered approach and protects agency and intellectual capacities in relation to uses of generative AI. The Recommendation on the Ethics of Neurotechnology (UNESCO, 2025), adopted by the 43rd General Conference in Samarkand on 11 November 2025, is a close institutional antecedent that requires precise comparison. Its scope is not limited to devices that interact directly with the nervous system: it also covers indirect neural data and non-neural data capable of enabling inferences about mental states—for example, eye tracking, keystroke dynamics, voice analysis, or facial emotion recognition—as well as the convergence of neurotechnology with artificial intelligence and other technologies (UNESCO, 2025, paras. 6 and 17). For priority analysis, its pre-adoption history is also material: UNESCO shared the first draft with Member States in September 2024, and an intergovernmental meeting of experts reviewed and approved the draft from 12 to 16 May 2025, before both cutoff dates used in this paper (UNESCO, 2024-2025). Paragraph 47 further extends protection against undue influence or manipulation to both the internal processing of thoughts and their external expression; and paragraph 122 requires, in education, assessment of risks such as dependency and loss of skills associated with the use of neurotechnology. These passages narrow the distance from the concern defended here and

should be expressly acknowledged. Their normative axis, however, remains access to, inference of, or modulation of mental states and, in education, the effects of neurotechnology use. They do not formulate as an autonomous object the non-manipulative, non-neural substitution of the human process of understanding by external generative systems. UNESCO's Recommendation on the Ethics of Artificial Intelligence (2021), for its part, treats autonomy and human oversight as general principles without configuring the cognitive process itself as the object of a right.

5.9. Regional Institutional Convergence: CEJA/JSCA and the Judicial Function

A particularly relevant regional convergence appears in the Justice Studies Center of the Americas (JSCA/CEJA). Its Spanish-language report *Inteligencia Artificial en los Poderes Judiciales: Reflexiones y Lineamientos para las Américas* [Artificial Intelligence in the Judiciary: Reflections and Guidelines for the Americas] (Pizarro-Zúñiga, 2025) distinguishes decision-support tools from uses involving adjudicative-content generation or substitution of judicial will; warns that human review may cease to be meaningful when excessive workloads prevent real and effective substantive oversight; states as an institutional objective that AI should not replace human rationality, preserving critical judgment, legal reflection, and substantive decision-making by persons; and systematizes human control, non-substitution of human rationality, auditability, and traceability among regional governance principles. It also notes that even decision-support uses may affect how facts are known and interpreted, how judicial conviction is formed, and how judgments are drafted.

This convergence is strong but sectoral. CEJA/JSCA addresses the judicial function, due process, judicial independence, and the preservation of human rationality or will; it does not formulate Cognitive Dignity as a general human right protecting the person's process of sense-making and knowledge formation. The official repository dates the Spanish edition only as December 2025, without a day. It is therefore not used here as evidence of priority over the 5 December 2025 cutoff, but as contemporaneous institutional convergence. An official English edition was published in May 2026 (Pizarro-Zúñiga, 2026).

5.10. Parallel Developments: The Literature on “Cognitive Sovereignty”

A number of recent works use the expression “cognitive sovereignty” for a neighboring problem. Brcic (2025), in a preprint predating the first cutoff date, defines it as the capacity of individuals, groups, and nations to maintain autonomous thought and preserve identity in the face of AI assistants with persistent memory; its framing is predominantly geopolitical and does not configure an individual right over one's process of understanding. After the cutoff dates, Konigsberg (2026) proposed cognitive sovereignty as an epistemological and psychological construct: the capacity to remain the genuine author of one's own understanding and to notice when one's thinking is being displaced. It is probably the contemporary formulation closest to the differential case defended here, although it lacks legal status. These works do not alter the chronology of the proposal, but they show that cognitive substitution is being identified in parallel across several disciplines, which reinforces its relevance and requires the originality hypothesis to remain strictly circumscribed.

5.11. Comparative Synthesis

Source	Main contribution	Relationship to Cognitive Dignity
Sententia (2004); Bublitz (2013)	Cognitive liberty and mental self-determination.	Partial legal antecedent: control over mind/technologies, not the complete formula of the process of sense-making in the face of substitution.
Ienca & Andorno (2017)	Cognitive liberty, mental privacy, mental integrity, psychological continuity.	Very close legal family, centered on neurotechnology, access/manipulation, and identity.
ICCPR Art. 18; McCarthy-Jones (2019); Alegre (2022); A/76/380 (2021)	Freedom of thought: forum internum, freedom from impermissible alteration, enabling environment.	Strongest legal competitor; evolutionary interpretation may absorb part of the core, but the doctrine centers on alteration and is absolute in character.
Cassinadri (2025)	Multidimensional and extended right to mental integrity.	Complementary regime: its test (non-consensual interference that bypasses reason and causes harm) does not reach consented, non-manipulative substitution.
Law 21,383 (Chile); IAJC (2023); A/HRC/57/61 (2024)	Positive law and regional/universal soft law on neurotechnology.	Neighboring field; its object is brain activity and neural data.
Boon (2021)	Human dignity in the face of substitution of capacities by AI.	Partial AI-dignity antecedent; not an autonomous processual right.
Teo (2023)	Human dignity, AI, and cognitive autonomy.	Strong partial antecedent; not the same legal object.

Source	Main contribution	Relationship to Cognitive Dignity
Bhakuni (2023)	Epistemic rights and dignity of the knower.	Strong partial antecedent; automation is not the core threat.
Faraoni (2023)	Mental self-determination against computational manipulation.	Closest legal antecedent; differs in core mechanism: manipulation/persuasion versus non-manipulative substitution.
UNESCO (2023)	Human agency and intellectual capacities in relation to GenAI.	Partial institutional antecedent to concern over cognitive substitution.
Muñoz & Marinaro (2024)	Habeas cogitationem: procedural safeguard for freedom of thought and neurorights.	Procedural instrument, not a substantive right; especially relevant to Habeas Mens (§11).
Malone et al. (2025)	Automation, dignity, and epistemic agency of experts.	Very close; does not formulate a general right to the inner process.
Coeckelbergh (2025)	AI and epistemic agency in belief formation/revision.	Describes a mechanism of impairment; not a Cognitive Dignity right.
Rughiniş et al. (2025)	Institutional legitimacy of human-AI configurations.	Important for attribution and responsibility; does not use dignity as the foundation of the human process of knowing.
Pizarro-Zúñiga / CEJA (2025)	Judicial AI: substantive human oversight, non-substitution of human rationality, auditability, and traceability.	Strong sectoral institutional convergence; not a general right to the cognitive process. Dated only “December 2025,” so it is not used to establish priority over 5 December.

Source	Main contribution	Relationship to Cognitive Dignity
Brcic (2025); Konigsberg (2026)	“Cognitive sovereignty” in relation to AI systems.	Brcic: geopolitical framing, predating the cutoffs. Konigsberg: construct very close to the differential case, postdating the cutoffs and lacking legal status.

Comparative result: the expression is not original, and there are strong partial and near antecedents. Within the audited corpus, no single source has been identified, prior to the cutoff dates, that simultaneously brings together: the person as knowing subject, the process of sense-making/knowledge formation as the immediate object, its own temporality and openness, technological substitution or foreclosure including non-manipulative forms, and the status of an emerging human right. This conclusion is necessarily provisional and falsifiable; it is not equivalent to proving universal non-existence.

6. What Does Cognitive Dignity Add to Existing Rights?

Cognitive Dignity does not begin from a “law-free zone.” The field is densely regulated by prior rights and principles. The legally serious question is whether a functional residue exists that those frameworks do not describe with sufficient clarity.

Right framework or	Primary coverage	Relevant residue
Human dignity	Non-instrumentalization and intrinsic value of the person.	Broad foundation; does not by itself determine when cognitive substitution becomes legally relevant.
Freedom of thought	Protection of the formation and maintenance of thought against impermissible interference.	Strong legal competitor (§5.3): evolutionary interpretation may absorb part of the core, especially through the enabling-environment route; its doctrine centers on impermissible alteration, and its absolute character complicates contextual balancing.

Right framework or	Primary coverage	Relevant residue
Cognitive liberty / mental self-determination	Control over mental states, experiences, and development.	Usually centers on access, alteration, or manipulation rather than external substitution of the process of understanding.
Mental integrity	Protection against harm or unauthorized interference in the mental dimension.	In Cassinadri's (2025) reconstruction, it requires non-consensual interference that bypasses reason and causes harm; it does not reach consented, non-manipulative substitution with gradual effects.
Mental privacy / data	Access, extraction, processing, and disclosure of mental data or content.	Does not cover cognitive substitution when no internal content is accessed.
Personal autonomy	Decisions and self-determination.	Formal choice may exist even when material understanding has been displaced.
Sector-specific rights	Due process, education, health, professional duties, oversight.	May resolve specific domains without constructing a general protection of the person as a knowing subject.

Right framework or	Primary coverage	Relevant residue
AI Act / Council of Europe Convention	Risks, human oversight (Art. 14, including automation bias, Art. 14(4)(b)), AI literacy (Art. 4), prohibition of manipulative techniques (Art. 5(1)(a)), rights, and governance.	Partial and heterogeneous coverage: Article 14 of the AI Act requires human oversight for high-risk systems and addresses automation bias; Article 4 imposes AI-literacy obligations with broader scope; and Article 5 establishes general prohibitions concerning certain manipulative techniques. These duties do not, however, configure a general right of the person over their own cognitive process, nor do they cover every case of non-manipulative cognitive substitution.

7. The Differential Case: Non-Manipulative Cognitive Substitution

The clearest incremental core appears when the problem is not influencing what a person thinks, but displacing the process by which they would have come to understand or work through that thought. Non-manipulative cognitive substitution may be defined as the material transfer to an external system of a function of understanding, evaluation, integration, or sense-making that, in the context at issue, is relevant to the person's status as a knowing agent, without any need directly to alter their beliefs or mental states.

For purposes of this article, “material” substitution should not be understood as a quantitative measure of how much a person has “thought” personally. It refers instead to the functional conditions of cognition in the context at issue. Relevant indicators include whether the person had a genuine opportunity to understand the relevant reasons, reconstruct or interrogate them, sufficient time to do so, a meaningful possibility of departing from the system's output, and a capacity to identify which cognitive functions were performed by the system and which remained attributable to the person. These indicators are contextual rather than cumulative requirements: their relevance depends on the function being performed, the consequences of the decision, the degree of asymmetry between the person and the system, and any applicable legal or institutional duty of personal judgment.

A minimal example helps show the difference. A person receives from a system a complete, correct, and transparent conclusion. No neural data are read; there is no deception, subliminal

persuasion, or coercion, and formally the person may accept or reject the result. Yet the environment demands speed, provides neither time nor information to reconstruct the reasons, and treats acceptance as if it implied the person's own understanding. Privacy and mental integrity may remain intact; decisional autonomy may formally exist. The question posed by Cognitive Dignity is different: was there a genuine opportunity for the person to act as the subject of the knowing process that the context attributes to them?

This case must be handled carefully to avoid paternalism. There is no general right to perform every calculation, search, or inference personally. People legitimately and beneficially delegate cognitive tasks. The problem arises when delegation ceases to be chosen or intelligible, when what was in fact displaced is presented as human understanding, or when the architecture materially prevents reconstruction, questioning, or appropriation of knowledge where such participation has substantive significance.

Throughout this article, references to AI systems “performing reasoning” refer to the functional production of outputs that perform, reproduce, or simulate reasoning-related operations. The expression does not presuppose that such systems possess human reasoning, understanding, or consciousness. The distinction is important because the normative question addressed here concerns the attribution and preservation of human cognitive processes, not the resolution of the philosophical question of whether machine outputs constitute reasoning in the human sense.

The empirical plausibility of the case does not depend on speculation. Human-factors literature documented decades ago that automating a task degrades the operator's ability to perform it when the system fails (Bainbridge, 1983), and that support systems generate automation bias and complacency even among expert users (Parasuraman & Manzey, 2010). Research on cognitive offloading shows that externalizing mental operations is a common and often rational strategy, but one that may carry costs for memory and one's own processing (Risko & Gilbert, 2016); recent studies associate intensive use of AI tools with lower critical-thinking performance mediated by cognitive offloading (Gerlich, 2025). These findings do not prove generalized harm—in Gerlich's case the evidence is correlational—but they are sufficient to reject the idea that cognitive substitution is an imaginary risk.

The empirical magnitude, persistence, and reversibility of AI-related cognitive substitution remain open questions. Existing evidence supports concern about automation bias, cognitive offloading, and changes in critical-thinking performance, but it does not establish a uniform causal effect across users, tasks, or institutional settings. Nor is it yet clear under which conditions any reduction in independent cognitive engagement is temporary, recoverable, or cumulative. The normative proposal advanced here therefore does not depend on demonstrating that cognitive delegation generally produces harm. Its narrower claim is that, where an institutional or technological architecture materially displaces a cognitively significant human function, the resulting risk to human status as a knowing subject may warrant legal safeguards even when the empirical trajectory and reversibility of the effect remain context-dependent.

7.1. Assistance, Delegation, and Substitution

Modality	Description	Prima facie compatibility
Assistance	Technology enhances or facilitates a capacity while the person retains sufficient understanding and control for the context.	Compatible.
Voluntary delegation	The person chooses to entrust a bounded cognitive task to a tool and reasonably understands what is being delegated.	Compatible, except in contexts involving non-delegable duties.
Material substitution	The architecture displaces the relevant cognitive function and leaves the person with merely formal approval or without genuine conditions for understanding/review.	Potentially problematic.
Manipulation	The architecture seeks to alter decisions, beliefs, or mental states through exploitation of vulnerabilities or covert persuasion.	A neighboring problem, already partly covered by other frameworks; it may also affect Cognitive Dignity.

7.2. A Contrast Case: Colombian Constitutional Court, Judgment T-323 of 2024

In 2024, the Second Review Chamber of the Colombian Constitutional Court reviewed a second-instance judgment in a tutela (constitutional protection) action in which the judge had used ChatGPT 3.5 to “extend the arguments” of a decision concerning the fundamental rights of a child with autism spectrum disorder. The Court examined whether the decision had been made by the judge or by the system, and whether it was adequately reasoned. It concluded that there had been no violation of due process because the tool was used after the judge had reasoned and adopted the decision; it accepted the use of AI as support for administrative, documentary, and summarization tasks under human oversight; it stated that AI may not replace the logical and human reasoning that belongs to each judge; and it ordered the Superior Council of the Judiciary to issue guidelines (Colombian Constitutional Court, 2024).

T-323 of 2024 should therefore be understood here as an illustrative case of the practical relevance of the underlying problem and of partial convergence with existing legal doctrines, not as judicial recognition of Cognitive Dignity as an autonomous right. The judgment demonstrates that existing law may already protect elements of the human cognitive process in AI-assisted decision-making; the distinct claim advanced here is that such protection may leave unresolved cases of non-manipulative substitution in which the decisive issue is not

merely when AI was used, but which cognitive functions were materially performed by the system and which remained genuinely attributable to the human decision-maker.

The case is valuable because it shows both the strength of existing frameworks and the precise location of the residue. On the facts decided, Cognitive Dignity and due process converge: the Court already protected the core defended here by requiring decision-making reasoning to remain human. That convergence must be acknowledged as evidence of partial absorption. Yet the operative criterion used—the temporal sequence, namely that the tool intervened after the decision—is a criterion of order and outcome, not of process. It does not answer what happens to the part of the reasoning drafted by the system that the legal order continues to attribute to the judge, nor cases in which consultation precedes the decision.

Consider a variation: the judge consults the system before deciding; the system proposes a correct, transparent, and non-manipulative solution together with its reasoning; under time pressure, the judge conducts a formal review and adopts it. The general principle stated by the Court—AI may not replace judicial reasoning—remains applicable, but it lacks an operational criterion: the rule of post-decision use no longer works, and human oversight may be treated as satisfied by nominal review. Cognitive Dignity asks questions here that change the diagnosis: whether there was a genuine opportunity to reconstruct the reasons (duty to provide a genuine opportunity), whether it is possible to reconstruct which functions were performed by the system and which by the judge (duty of functional traceability), and whether something substantively displaced is presented as human judgment (duty to prevent concealed substitution). It also changes the remedy: disclosure of the tool's use is not enough; the human cognitive contribution must be documented. This shift in criterion—from sequence to process—is the incrementality defended by this paper, and it is also the point at which Habeas Mens (§11) supplies the evidentiary standard.

8. Gap and Incrementality Test

The creation of a new right cannot be justified merely by the appearance of a technology. To assess it, this paper adopts, with adaptations, the quality-control criteria proposed by Alston (1984) against the proliferation of new human rights—fundamental social value, not merely duplicating an existing right, sufficient precision to generate identifiable rights and obligations, and capacity to achieve international consensus—and Nickel's (2007) criteria of a common and serious threat and justifiable burdens on duty-bearers. It adds a further, more demanding criterion: the need for legal autonomy, meaning that the protection cannot be achieved through doctrinal development of already recognized rights. Applied to Cognitive Dignity, the result is not a total gap but a residual one.

Criterion	Result	Basis
Fundamental human value	Passes	The possibility of understanding and making sense is linked to dignity, autonomy, and freedom of thought.

Criterion	Result	Basis
Common and serious threat	Passes with reservation	Overdependence and automation of cognitive tasks are risks with empirical support (§7); their magnitude depends on context, and part of the evidence is correlational.
Non-redundancy	Partially passes	A residual case of non-manipulative cognitive substitution remains, although existing rights might absorb it through evolutionary interpretation.
Precision	Passes after revision	It must be limited to materially relevant substitution or interference; not every form of assistance, acceleration, or influence.
Proportionate burdens	Passes after revision	The proposed obligations can be modulated by risk, consent, vulnerability, and function.
International consensus	Does not pass at present	There is no international recognition of an autonomous right called Cognitive Dignity.
Need for legal autonomy	Not demonstrated	Dignity, freedom of thought, autonomy, or cognitive liberty could accommodate the protection through doctrinal development.

Material incrementality may therefore be regarded as defensible; the necessary legal autonomy of a separate right may not. This distinction is central. A concept may contribute a useful normative description and a cross-cutting architecture without requiring the creation of an independent human right.

9. Proposed Minimum Normative Content

If Cognitive Dignity is to be legally operational, it must be translated into verifiable duties rather than an abstract prohibition on “intervening” in the mind. A minimum content could include the following rules:

- Duty to prevent concealed substitution: do not present as human understanding, evaluation, or judgment a result whose substantive cognitive function was displaced to a system when human involvement is legally relevant.
- Duty of context-specific intelligibility: provide sufficient information and reasons for the person to understand what is necessary in the context, without converting this into an unlimited right to technical explanations.

- Duty to provide a genuine opportunity: provide reasonable time and conditions to review, question, or depart from an automated result in contexts involving significant consequences.
- Duty of functional traceability: make it possible to reconstruct which functions were performed by the system and which were actually performed by the person when cognitive attribution has legal or institutional effects.
- Duty of cognitive non-coercion: prevent a formal possibility of departing from the system from being emptied of substance by tempos, penalties, or burdens that make independent judgment unworkable.
- Duty of heightened protection in contexts of severe asymmetry or gravity, such as justice, health, education, or public administration.
- Assistance and accessibility exception: a tool that enhances capacities, compensates for a disability, or provides chosen support does not infringe the right merely because it performs part of the cognitive work.
- The addressee of these duties is necessarily context-dependent. Depending on the setting, relevant duty-bearers may include the State and public authorities, courts and adjudicative bodies, educational institutions, employers and professional organizations, and the designers, deployers, or operators of technological systems. The proposal does not assume that all such actors bear identical obligations. Rather, the applicable duty depends on who structures, controls, or institutionally attributes the relevant cognitive function, and on the legal or professional consequences attached to that function. In some contexts the primary obligation will concern the institution that organizes the decision process; in others, it may concern the human decision-maker or the provider or deployer of the technological architecture. The purpose of this distinction is to identify the institutional location of the relevant safeguard without presupposing a complete regime of liability.

These obligations should not be confused with the full historical content of the concept. They constitute a proposed legal-operational translation intended to make the concept applicable without losing its original object: the person as subject of their own process of knowing.

10. Areas of Application

10.1. Education

In education, the problem arises when a system produces answers before sufficient independent elaboration has occurred and the educational activity continues to be assessed as if it reflected the student's understanding. Cognitive Dignity does not require banning generative assistants; it requires distinguishing when a tool enhances learning and when it materially replaces it. The legal variable is not the use of AI as such, but the relationship among pedagogical function, cognitive participation, and attribution of the result.

10.2. Health

In medicine, support systems can improve accuracy and reduce errors. The proposal does not require every professional to reproduce internally all of the system's calculations. It does, however, raise a question when a decision is professionally attributed to a person who lacks real conditions to understand the relevant reasons, challenge the output, or recognize its

limits. Here Cognitive Dignity intersects with professional duties, consent, safety, and clinical responsibility.

10.3. Justice and Public Administration

In governmental or judicial decisions, the issue may be even more demanding because the legal order assigns certain functions to a human authority. Cognitive Dignity provides the general foundation for asking whether the human process has been preserved; doctrines of reason-giving, due process, non-delegation, and meaningful human control provide sector-specific rules. When the problem becomes specifically one of proving who performed the decision-making cognitive function and whether the judgment is attributable to the human being, Habeas Mens enters as a distinct later development. The Colombian case analyzed in §7.2 illustrates this articulation.

10.4. Work and Expert Professions

Automation may formally keep a professional “in the loop” while reducing their real capacity to depart from the system, review reasons, or exercise expert judgment. The epistemic agency studied by Malone et al. illustrates this risk. Cognitive Dignity adds a general formulation centered on the status of the knowing subject, but its application must respect legitimate organization of work and avoid turning every technological dependency into a human-rights violation.

11. Cognitive Dignity and Habeas Mens

Cognitive Dignity and Habeas Mens share a genealogy but perform different functions. The former is the general proposal to protect the person as the subject of the process of sense-making and knowledge formation. Habeas Mens appears later as a legal-operational development aimed at decisions attributed to a person: it seeks to verify whether the cognitive function of forming judgment was actually performed by the human decision-maker and with what traceability. (Sinopoli, 2026a).

The distinction avoids two confusions. First, Cognitive Dignity is not reducible to meaningful human oversight or non-delegation; its object is broader than the institutional attribution of a decision. Second, Habeas Mens should not be used retrospectively to redefine the 2025 concept. It may be a sector-specific and evidentiary application of a problem of cognitive integrity, but it is neither the concept's original source nor its synonym. (Sinopoli, 2025d; 2026a).

CEJA/JSCA's 2025 judiciary report is especially close to this interface: it calls for human control, non-substitution of human rationality, auditability, traceability, and substantive rather than merely formal oversight. That institutional convergence does not amount to Habeas Mens, because CEJA/JSCA does not articulate an evidentiary test for attributing the cognitive-decisional function to the human decision-maker; it nevertheless shows that the problem operationalized by Habeas Mens has become institutionally salient in regional judicial governance (Pizarro-Zúñiga, 2025).

The distinction also has a priority dimension that should be made explicit. The expression “habeas mens” has an earlier nominal use in Sommaggio, Mazzocca, Gerola, and Ferro (2017), as a programmatic formula without an operational standard. More relevant is Muñoz

and Marinaro's (2024) proposal for habeas cogitationem: a procedural remedy, analogous to habeas corpus and habeas data, intended to make freedom of thought and, subsidiarily, neurorights effective. The proposal has received institutional attention, including Argentine Bill No. 2815-D-2025, introduced on 30 May 2025 by Deputy Ramiro Gutiérrez and entitled “Régimen de protección del fuero interno de la libertad de pensamiento y los neuroderechos afines. Habeas Cogitationem” [Regime for the Protection of the Internal Forum of Freedom of Thought and Related Neurorights. Habeas Cogitationem]. Habeas Mens should therefore not be presented as the first procedural safeguard of the mind: its difference lies in its object—the human attributability of judgment in AI-assisted decisions, rather than protection against neurotechnological intrusions—and in its character as an evidentiary standard rather than an amparo-type protective action.

12. Main Objections

12.1. Rights Inflation

The strongest objection is that not every valuable interest requires a new human right. Alston (1984) warned early against the creation of rights without quality control, and the neurorights debate has reproduced this tension: Bublitz (2022) argued that much of the debate suffered from conceptual problems before reconstructing its defensible core; Hertz (2023) argued that freedom of thought, properly reinterpreted, renders several neurorights unnecessary; and Muñoz and Marinaro (2023) reconstructed neurorights as reconceptualizations of existing rights. Freedom of thought, dignity, autonomy, and cognitive liberty could evolve in the same way to cover the cases analyzed here. This objection is valid and, at present, prevents any claim that an autonomous right is legally indispensable. The proposal may function as an emerging doctrine even without independent positivization.

12.2. Cognitive Paternalism

Protecting the human process must not become an obligation to think “without assistance.” People also have the right to delegate, simplify, rely on experts, and use tools. The threshold should therefore focus on material substitution that is non-consensual, opaque, or functionally decisive, especially when the architecture attributes to the subject an understanding that the subject had no genuine opportunity to develop.

12.3. Vagueness of the “Inner Process”

The historical formulation used broad symbolic and process-oriented language. For legal use, it must be translated into observable variables: opportunity for understanding, possibility of reconstructing reasons, capacity to question, sufficient time, genuine freedom to depart from the system, functional traceability, and attribution of the result. This translation does not eliminate the original concept; it makes it operational.

12.4. Independence from the Theoretical Framework of Origin

Cognitive Dignity arose within a broader theoretical program developed by the author—Comparative Symbolic Epistemology, with the TESA protocol and the OCD construct (Sinopoli, 2025b)—which includes hypotheses concerning a grammar of the symbolic process and the role of REM sleep. Its legal validity should not depend on those hypotheses. They may retain historical or research value, but a universal right must be capable of

justification through recognizable human interests and normative risks independently of any particular psychological theory.

12.5. Originality and Falsification

The originality claim must remain falsifiable. If a source predating the cutoff dates is found that brings together in a single construction a right of the person as knowing subject, the process of sense-making or knowledge formation as its object, technological substitution or foreclosure as the threat, and an equivalent normative status, the specific priority claim must be revised. Until such a source appears, the prudent formulation is that no equivalent individual source has been identified within the audited corpus. (Sinopoli, 2026b; see Annex A).

13. Current Legal Status of the Proposal

Question	Answer
Is the expression “Cognitive Dignity” original?	No. Earlier uses exist with different meanings.
Is concern with mind, dignity, or AI original?	No. There are abundant antecedents.
Are there close conceptual antecedents?	Yes. Especially freedom of thought (A/76/380, McCarthy-Jones, Alegre), cognitive liberty and neurorights, mental integrity (Cassinadri), Teo, Bhakuni, Faraoni, UNESCO, Malone, Coeckelbergh, and the literature on cognitive sovereignty.
Is there an equivalent earlier individual antecedent in the audited corpus?	None identified.
Does specific originality of the formula survive?	Yes, as a hypothesis supported within the audited corpus, not as demonstrated universal priority.
Is there a total legal gap?	No.
Is there a residual gap?	Yes, especially in cases of non-manipulative cognitive substitution.
Has the need for an autonomous human right been demonstrated?	No.

Question	Answer
Can it be presented as a proposed emerging human right?	Yes, with limits and doctrinal revision, and without confusing the proposal with positive law in force.

The defensible contribution, therefore, does not lie in having discovered that technology can affect the mind. It lies in proposing the human process of sense-making and knowledge formation itself—including its non-substitution by technology—as a specific object of legal protection, and in organizing under the category of Cognitive Dignity a problem that existing frameworks cover only partially, diffusely, or sectorally.

14. Conclusion

Cognitive Dignity poses a legal question that becomes increasingly relevant as artificial intelligence ceases to be merely an information tool and begins to perform functions of understanding, synthesis, evaluation, and the production of reasons. The problem is not only what data are processed, what thought is manipulated, or who formally adopts the decision. It also matters whether the person retains genuine conditions to exist as the subject of the knowing process that the surrounding context attributes to them.

The comparative inquiry requires a double moderation. On the one hand, the proposal cannot be presented as the creation of the expression “cognitive dignity,” nor as the first concern with mental autonomy, human dignity in the face of AI, or epistemic agency. Those antecedents exist and are substantive. On the other hand, the proposal should not be reduced to a mere sum of those pieces: within the audited corpus, no earlier individual source has been identified that formulates the same complete architecture, and Faraoni (2023), although very close, is structured around manipulation and mental self-determination rather than non-manipulative cognitive substitution. Freedom of thought, evolutionarily reinterpreted, remains the most serious competitor, and the Colombian case shows that, in cases already resolved, existing frameworks and Cognitive Dignity may converge.

The most defensible doctrinal position is therefore intermediate: Cognitive Dignity survives as a proposal for an emerging human right with an identifiable differential core, but not as an already recognized autonomous right or as a demonstrated necessity for legal autonomy. Its future will depend on two tests: showing cases in which it reasonably changes the diagnosis, standard, or remedy relative to existing rights, and translating its process-oriented intuitions into verifiable thresholds capable of distinguishing legitimate assistance, voluntary delegation, and materially degrading substitution.

In that sense, the proposal does not close the debate on freedom of thought, autonomy, or neurorights. It seeks to offer an additional category for a specific transformation of the age of artificial intelligence: the possibility that a human being may remain present in a decision or knowledge process while ceasing to be, in any meaningful sense, the subject who understands, works through, and makes sense.

Transparency Note on Artificial-Intelligence Tools

Artificial-intelligence tools were used as technical support in the preparation, comparison, and editing of this article. Selection of the problem, documentary genealogy, formulation of Cognitive Dignity, conceptual decisions, and final responsibility for the content belong to the author. AI tools are not presented as co-authors.

Annex A. Corpus Search and Audit Protocol

The claims in this paper concerning the absence of an equivalent individual antecedent are limited to the corpus described in this annex. Its purpose is to allow any reader to reproduce or expand the search and, potentially, falsify the hypothesis of specific originality (§12.5).

Sources consulted. A directed and adversarial doctrinal search was conducted, not a systematic PRISMA-type bibliometric review. Candidate identification combined structured web searches with consultation and verification in SSRN, PubMed/PMC, SpringerLink, Frontiers, BMJ Global Health, arXiv, university repositories (including University of Twente Research Information), and institutional or official sources of UNESCO, the United Nations, the OAS/Inter-American Juridical Committee, the European Union/EUR-Lex, the Council of Europe, the Colombian Constitutional Court, the Supreme Court of Chile, the Chamber of Deputies of the Argentine Nation, and CEJA/JSCA. Where a primary source, publisher's page, or institutional repository was available, that version was preferred over secondary summaries.

Languages. The core searches were conducted in Spanish and English. One Dutch-language source (Boon, 2021) was also included on the basis of its institutional record, metadata, and available text. No exhaustive search was conducted in all languages; this limitation precludes any claim of absolute universal priority.

Search strings. Literal strings and combinations included: “cognitive dignity”; “dignidad cognitiva”; “cognitive dignity” AND AI; “cognitive substitution” AND “human rights”; “freedom of thought” AND “artificial intelligence”; “freedom of thought” AND “enabling environment”; “cognitive liberty” AND technology; “mental integrity” AND AI; “mental self-determination” AND AI; “epistemic agency” AND automation; “human dignity” AND artificial intelligence; “cognitive sovereignty”; “cognitive offloading” AND AI; “habeas cogitationem”; “neurorights” AND freedom of thought; “judicial AI” AND “human oversight”; “no sustitución de la racionalidad humana” [“non-substitution of human rationality”]; and specific searches by author, title, DOI, legislative file number, or international-document number when a relevant candidate emerged.

Temporal window. No rigid lower bound was fixed: conceptual and legal antecedents were traced retrospectively to sources predating the contemporary neurorights literature whenever relevant. Two cutoffs were used for priority: 21 October 2025 for the epistemic-protective core and 5 December 2025 for the explicit legal formulation. Sources after 5 December 2025 were used only to identify parallel developments, convergence, or subsequent evolution, through 24 September 2026, and not to challenge priority retrospectively.

Inclusion and exclusion criteria. Included were: (a) sources predating the cutoffs that overlapped in one or more relevant dimensions—protected subject, processual object, threat, mechanism of harm, normative status, or correlative duty; (b) international instruments,

legislation, case law, and soft law protecting the mind, cognition, or human oversight; (c) doctrinal works on freedom of thought, cognitive liberty, mental integrity, mental self-determination, dignity of the knower, epistemic agency, automation, and cognitive sovereignty; and (d) especially close post-cutoff developments, solely as convergence. Excluded were duplicates, secondary summaries when the primary source was available, AI-ethics texts lacking a material connection to the cognitive process, and purely lexical coincidences without a comparable conceptual function. Terminological antecedents were retained only to rule out priority over the expression itself.

Selection flow. The search was iterative: identification through terms and cross-references; initial screening by title, abstract, and metadata; reading of the full text or decisive section when a candidate showed proximity; and dimension-by-dimension comparison. Each source was classified as conceptual anticipation, near anticipation, partial antecedent, or analogy/terminological antecedent. No artificial number of “records identified” is reported because the search did not begin from a single closed query or a bibliometric database with a stable universe. The final set of sources reaching comparative relevance is made explicit in §5.11 and the bibliography, allowing the reader to examine each candidate and propose others.

Use of AI tools in the search. ChatGPT (OpenAI) and Claude (Anthropic) were used as auxiliary tools to generate search terms, locate possible antecedents, perform adversarial comparison, and detect bibliographic errors or omissions. No AI output was accepted as evidence in itself. References and decisive claims incorporated into the paper were checked, when the source was available, against the primary text, the institution's official page, the academic publisher, or an institutional repository; where that was not possible, a reliable bibliographic record was used, and the source was not attributed more than that record permits. Conceptual classification and final responsibility for the conclusion belong to the author.

Limitations. The search cannot exclude non-indexed sources, non-digitized gray literature, or texts in languages not covered; accordingly, the conclusion in §5.11 is provisional and falsifiable.

Availability of primary sources. The author's manuscripts cited as Sinopoli (2025a-e, 2026a-b) are registered with Safe Creative, and the references provide proof number, date, and time where applicable. At the time of this version, not all have a DOI or an independent public repository. This paper therefore reproduces the decisive formulations and their documentary-fixation data; a later open deposit of the manuscripts or a documentary annex would further facilitate external verification, but no such deposit is presumed to have occurred in this version.

References

Primary Sources by the Author

1. Sinopoli, A. (2025a). *Epistemología Simbólica Comparativa* [Comparative Symbolic Epistemology]. Unpublished manuscript. Safe Creative, technological proof 2509173099388, 17 September 2025, 15:56 UTC.

2. Sinopoli, A. (2025b). Epistemología Simbólica Comparativa (ESC), incluyendo el Protocolo TESA y el Constructo OCD: Marco Epistemológico General para el Análisis de Estructuras Simbólicas Activas [Comparative Symbolic Epistemology (ESC), including the TESA Protocol and the OCD Construct: General Epistemological Framework for the Analysis of Active Symbolic Structures]. Unpublished manuscript. Safe Creative, technological proof 2510153326103, 15 October 2025, 15:56 UTC.
3. Sinopoli, A. (2025c). ESC TESA OCD v2.3 - Epistemología Simbólica y Optimización Ética del Conocimiento Humano [Symbolic Epistemology and Ethical Optimization of Human Knowledge]. Unpublished manuscript. Safe Creative, technological proof 2510213392123, 21 October 2025, 22:19 UTC.
4. Sinopoli, A. (2025d). EL TRATADO - EPISTEMÍA: Hacia un derecho humano emergente: la protección del espacio simbólico en la era digital [THE TREATISE - EPISTEMÍA: Toward an Emerging Human Right: Protecting the Symbolic Space in the Digital Age]. Unpublished manuscript. Safe Creative, technological proof 2512053942261, 5 December 2025, 13:59 UTC.
5. Sinopoli, A. (2025e). EL TRATADO - EPISTEMÍA: Hacia un derecho humano emergente: la protección del espacio simbólico humano [THE TREATISE - EPISTEMÍA: Toward an Emerging Human Right: Protecting Human Symbolic Space]. Unpublished manuscript. Safe Creative, technological proof 2512204078191, 20 December 2025, 15:24 UTC.
6. Sinopoli, A. (2026a). Habeas Mens. Unpublished manuscript. Safe Creative, technological proof 2605065552815, 6 May 2026, 22:57 UTC.
7. Sinopoli, A. (2026b). Expediente de Auditoría - Dignidad Cognitiva [Audit Dossier - Cognitive Dignity]. Final updated closure, 24 September 2026. Research manuscript.

Bibliography and External Sources

8. Alegre, S. (2022). Freedom to Think: The Long Struggle to Liberate Our Minds. Atlantic Books.
9. Alston, P. (1984). Conjuring Up New Human Rights: A Proposal for Quality Control. *American Journal of International Law*, 78(3), 607-621.
<https://doi.org/10.2307/2202599>
10. Argentina, Chamber of Deputies of the Nation. (2025). Bill 2815-D-2025, “Régimen de protección del fuero interno de la libertad de pensamiento y los neuroderechos afines. Habeas Cogitationem” [Regime for the Protection of the Internal Forum of Freedom of Thought and Related Neurorights. Habeas Cogitationem]. Introduced by Ramiro Gutiérrez on 30 May 2025, *Trámite Parlamentario* No. 64.
<https://diputados.gob.ar/diputados/crgutierrez/proyecto.html?exp=2815-D-2025>
11. Bainbridge, L. (1983). Ironies of Automation. *Automatica*, 19(6), 775-779.
[https://doi.org/10.1016/0005-1098\(83\)90046-8](https://doi.org/10.1016/0005-1098(83)90046-8)
12. Bhakuni, H. (2023). Epistemic repair in global health: a human rights approach towards epistemic justice [Commentary]. *BMJ Global Health*, 8(8), e013544.
<https://doi.org/10.1136/bmjgh-2023-013544>
13. Boon, M. (2021). Artificiële Intelligentie in de Samenleving, en de vraag naar menselijke waardigheid [Artificial Intelligence in Society, and the Question of Human Dignity]. Professional e-book associated with the Reflecties exhibition and the Man and Machine – Learning in the Digital Society project, University of Twente/Tetem.

- University of Twente Research Information, published 29 April 2021.
<https://research.utwente.nl/en/publications/artifici%C3%ABle-intelligentie-in-de-samenleving-en-de-vraag-naar-mens/>
14. Brcic, M. (2025). The Memory Wars: AI Memory, Network Effects, and the Geopolitics of Cognitive Sovereignty [preprint]. arXiv:2508.05867.
 15. Bublitz, J. C. (2013). My Mind Is Mine!? Cognitive Liberty as a Legal Concept. In E. Hildt & A. G. Franke (eds.), *Cognitive Enhancement: An Interdisciplinary Perspective* (pp. 233-264). Springer.
 16. Bublitz, J. C. (2022). Novel Neurorights: From Nonsense to Substance. *Neuroethics*, 15, 7. <https://doi.org/10.1007/s12152-022-09481-3>
 17. Cassinadri, G. (2025). The Right to Mental Integrity: Multidimensional, Multilayered and Extended. *Neuroethics*, 18, 16. <https://doi.org/10.1007/s12152-025-09585-6>
 18. Chile. (2021). Law 21,383, amending the Constitution to establish scientific and technological development at the service of persons. *Official Gazette*, 25 October 2021.
 19. Chile, Supreme Court (Third Chamber). (2023). *Girardi Lavín v. Emotiv Inc.*, Case No. 105.065-2023, judgment of 9 August 2023.
 20. Coeckelbergh, M. (2025). AI and Epistemic Agency: How AI Influences Belief Revision and Its Normative Implications. *Social Epistemology*, 40(1), 59-71. <https://doi.org/10.1080/02691728.2025.2466164>
 21. Colombia, Constitutional Court (Second Review Chamber). (2024). Judgment T-323 of 2024, 2 August 2024.
 22. Council of Europe. (2024). Framework Convention on Artificial Intelligence and Human Rights, Democracy and the Rule of Law, CETS No. 225. Opened for signature on 5 September 2024.
 23. European Union. (2024). Regulation (EU) 2024/1689 laying down harmonised rules on artificial intelligence (Artificial Intelligence Act). *Official Journal of the European Union*, L, 12 July 2024.
 24. Faraoni, S. (2023). Persuasive Technology and computational manipulation: hypernudging out of mental self-determination. *Frontiers in Artificial Intelligence*, 6, 1216340. <https://doi.org/10.3389/frai.2023.1216340>
 25. Fornet-Betancourt, R. (2007). La filosofía intercultural desde una perspectiva latinoamericana [Intercultural Philosophy from a Latin American Perspective]. *Solar*, 3(3), 23-40.
 26. Gerlich, M. (2025). AI Tools in Society: Impacts on Cognitive Offloading and the Future of Critical Thinking. *Societies*, 15(1), 6. <https://doi.org/10.3390/soc15010006>
 27. Hertz, N. (2023). Neurorights – Do We Need New Human Rights? A Reconsideration of the Right to Freedom of Thought. *Neuroethics*, 16, 5. <https://doi.org/10.1007/s12152-022-09511-0>
 28. Ienca, M., & Andorno, R. (2017). Towards new human rights in the age of neuroscience and neurotechnology. *Life Sciences, Society and Policy*, 13, 5. <https://doi.org/10.1186/s40504-017-0050-1>
 29. Institut de Drets Humans de Catalunya. (2009). Declaración Universal de Derechos Humanos Emergentes [Universal Declaration of Emerging Human Rights]. Final text approved at the Universal Forum of Cultures, Monterrey, 2007. Barcelona: Institut de

- Drets Humans de Catalunya. <https://www.idhc.org/es/publicaciones/declaracion-universal-de-derechos-humanos-emergentes/>
30. Inter-American Juridical Committee (OAS). (2023). Inter-American Declaration of Principles regarding Neuroscience, Neurotechnologies, and Human Rights. CJI/RES. 281 (CII-O/23) corr.1, 9 March 2023. https://www.oas.org/en/sla/iajc/themes_recently_concluded_Neuroscience_neurotechnologies_and_human_rights.asp
 31. Konigsberg, A. (2026). Cognitive Sovereignty: The Authorship Problem in AI-Assisted Thought. SSRN Working Paper, posted 20 April 2026. SSRN 6575778. <https://doi.org/10.2139/ssrn.6575778>
 32. Malone, E., Afroogh, S., D'Cruz, J., & Varshney, K. R. (2025). When Trust is Zero Sum: Automation's Threat to Epistemic Agency. *Ethics and Information Technology*, 27(2), 29. <https://doi.org/10.1007/s10676-025-09838-1> (preprint: arXiv:2408.08846, 2024).
 33. McCarthy-Jones, S. (2019). The Autonomous Mind: The Right to Freedom of Thought in the Twenty-First Century. *Frontiers in Artificial Intelligence*, 2, 19. <https://doi.org/10.3389/frai.2019.00019>
 34. Muñoz, J. M., & Marinaro, J. Á. (2023). Neurorights as reconceptualized human rights. *Frontiers in Political Science*, 5, 1322922. <https://doi.org/10.3389/fpos.2023.1322922>
 35. Muñoz, J. M., & Marinaro, J. Á. (2024). "You shall have the thought": habeas cogitationem as a New Legal Remedy to Enforce Freedom of Thinking and Neurorights. *Neuroethics*, 17, 18. <https://doi.org/10.1007/s12152-024-09551-8>
 36. Nickel, J. W. (2007). *Making Sense of Human Rights* (2nd ed.). Blackwell.
 37. Parasuraman, R., & Manzey, D. H. (2010). Complacency and Bias in Human Use of Automation: An Attentional Integration. *Human Factors*, 52(3), 381-410. <https://doi.org/10.1177/0018720810376055>
 38. Pizarro-Zúñiga, P. (2025). *Inteligencia Artificial en los Poderes Judiciales: Reflexiones y Lineamientos para las Américas [Artificial Intelligence in the Judiciary: Reflections and Guidelines for the Americas]*. Technical report. Centro de Estudios de Justicia de las Américas (CEJA/JSCA), December 2025. ISBN 978-956-6413-06-6. <https://biblioteca.cejamericas.org/handle/2015/5749>
 39. Pizarro-Zúñiga, P. (2026). *Artificial Intelligence in the Judiciary: Reflections and Guidelines for the Americas*. Justice Studies Center of the Americas (JSCA/CEJA), May 2026. ISBN 978-956-6413-09-7. <https://biblioteca.cejamericas.org/handle/2015/5752>
 40. Risko, E. F., & Gilbert, S. J. (2016). Cognitive Offloading. *Trends in Cognitive Sciences*, 20(9), 676-688. <https://doi.org/10.1016/j.tics.2016.07.002>
 41. Rughiniș, C., Vulpe, S.-N., Țurcanu, D., & Rughiniș, R. (2025). AI at the knowledge gates: institutional policies and hybrid configurations in universities and publishers. *Frontiers in Computer Science*, 7, 1608276. <https://doi.org/10.3389/fcomp.2025.1608276>
 42. Sententia, W. (2004). Neuroethical Considerations: Cognitive Liberty and Converging Technologies for Improving Human Cognition. *Annals of the New York Academy of Sciences*, 1013, 221-228. <https://doi.org/10.1196/annals.1305.014>

43. Sommaggio, P., Mazzocca, M., Gerola, A., & Ferro, F. (2017). Cognitive liberty. A first step towards a human neuro-rights declaration. *BioLaw Journal – Rivista di BioDiritto*, (3), 27-45. <https://doi.org/10.15168/2284-4503-255>
44. Sorum, P. C. (2017). In Search of Cognitive Dignity: The Diagnostic Challenges of Primary Care. *Medical Decision Making*, 37(1), 6-8. <https://doi.org/10.1177/0272989X16662643>
45. Teo, S. A. (2023). Human Dignity and Artificial Intelligence: Mapping the Contours and Utility of Human Dignity in Addressing Challenges Presented by AI. *Law, Innovation and Technology*, 15(1), 241-279. <https://doi.org/10.1080/17579961.2023.2184132>
46. UNESCO. (2021). Recommendation on the Ethics of Artificial Intelligence.
47. UNESCO. (2023). Guidance for Generative AI in Education and Research.
48. UNESCO. (2024-2025). Towards an International Instrument: Recommendation on the Ethics of Neurotechnology. Development history: first draft shared with Member States in September 2024; intergovernmental meeting of experts, 12-16 May 2025. <https://www.unesco.org/en/ethics-neurotech/recommendation>
49. UNESCO. (2025). Recommendation on the Ethics of Neurotechnology. Adopted by the 43rd General Conference in Samarkand on 11 November 2025. <https://www.unesco.org/en/legal-affairs/recommendation-ethics-neurotechnology>
50. United Nations, Human Rights Council Advisory Committee. (2024). Impact, opportunities and challenges of neurotechnology with regard to the promotion and protection of all human rights. A/HRC/57/61.
51. United Nations, Special Rapporteur on freedom of religion or belief (Ahmed Shaheed). (2021). Interim report: Freedom of thought. A/76/380.