

**FROM EPISTEMIC RISK TO RESPONSIBILITY ARCHITECTURE:
AI, HUMAN JUDGMENT, AND THE INSTRUMENTUM VOCALE
DOCTRINE**

Kildeev Adel

Leningrad State University named after A.A. Zhdanov (1990)

PhD in Law

ORCID: 0009-0001-2211-4944

Abstract: This article argues that governance of generative AI must move beyond hallucination, bias, explainability, alignment, and reliability toward an architecture of human responsibility. Drawing on Yang et al.'s analysis of persuasion, cognitive offloading, and epistemic lock-in, as well as Bednar et al.'s study of AI-assisted legal reasoning, it shows that generative systems may affect not only individual outputs but also the conditions under which users verify claims, exercise judgment, and preserve cognitive independence.

These findings are interpreted through the *Instrumentum Vocale* doctrine. Large language models should be understood not as epistemic agents, artificial experts, or legal actors, but as *instrumenta vocalia*: probabilistic speaking instruments that may assist drafting, organization, retrieval, and preliminary synthesis without possessing understanding, voluntas, judgment, or responsibility.

The article argues that retrieval, self-critique, prompting, multi-model comparison, and apparent auditability may improve practical reliability but cannot create an accountable bearer of judgment. Responsibility remains non-delegable and attached to the human actor or institution that deploys, verifies, and acts upon AI-generated output.

The article concludes that responsible AI governance requires a visible human chain of supervision, source verification, attribution, and final judgment. Such an architecture can reduce both catastrophic risks arising from hallucinated or unverified outputs and the gradual degradation of human professional and cognitive capacity through epistemic dependence.

Keywords: Generative AI; large language models; epistemic risk; cognitive offloading; epistemic lock-in; verification; human judgment; responsibility architecture; Instrumentum Vocale doctrine; AI governance.

The use of generative models ChatGPT (OpenAI), Gemini (Google), Grok (xAI) in drafting this article illustrates, rather than undermines, its central thesis: LLMs function as *instrumenta vocalia*—tools assisting in linguistic production without possessing understanding, intention, or authorship. Their involvement does not alter the locus of legal or intellectual responsibility, which remains exclusively with the human author, who makes all substantive, structural, and evaluative decisions. LLM outputs

cannot be accorded independent legal or intellectual status apart from the human author. These systems were used exclusively for structural outlining, comparative review, grammatical refinement, stylistic polishing, and the generation of draft formulations subject to human selection and revision. They were not used as sources of factual authority, legal verification, theological interpretation, mathematical proof, or independent scholarly judgment. All substantive claims, citations, structural decisions, and final evaluative conclusions remain the responsibility of the human author.

Chapter I. Introduction: From Epistemic Risk to Responsibility Architecture

The current debate on generative artificial intelligence is dominated by a familiar vocabulary: hallucination, bias, opacity, explainability, alignment, safety, and reliability. These concepts are important, but they do not yet reach the central institutional problem. A probabilistic system may become more accurate, more disciplined, and more useful while remaining incapable of understanding, judgment, intention, or responsibility. The question is therefore not simply whether a model produces fewer errors. The more fundamental question is what legal, professional, and epistemic status should be assigned to an output that may be linguistically persuasive but is not the product of an accountable mind.

Recent empirical research increasingly demonstrates that the use of generative AI can alter not only the speed of work, but also the quality of human reasoning, verification practices, and professional judgment. These findings should not be reduced to a simplistic claim that AI is inherently harmful or that human users are inevitably passive. The deeper concern is structural. Where a system produces fluent answers, citations, summaries, analyses, or recommendations at negligible marginal cost, institutions may gradually reorganize themselves around the availability of output rather than around the conditions of judgment. Verification becomes optional; expertise becomes simulated; responsibility becomes diffused across software providers, users, supervisors, and opaque technical systems.

This risk is particularly acute in law, medicine, engineering, education, public administration, and other domains where language is not merely communicative, but operative. A legal memorandum can influence litigation strategy. A medical summary can shape treatment. An engineering assessment can affect physical safety. In each of these settings, the relevant issue is not whether a model can assist a qualified professional. It plainly can. The issue is whether assistance is silently converted into delegation, and whether delegation is then mistaken for the transfer of responsibility.

The distinction matters because generative systems do not verify the truth of their own outputs in the human sense. They generate statistically plausible linguistic continuations. Retrieval tools, citations, system prompts, multi-model comparison, and alignment layers may improve practical reliability; they may also reduce certain types of error. But none of these mechanisms creates a bearer of judgment. A citation is not verification. Retrieval is not evaluation. Repetition across several models is not independent confirmation. And an apparently cautious system remains a system without voluntas, professional duty, or the capacity to answer for consequences.

This article is not another catalogue of AI failures. It does not attempt to reproduce the growing evidentiary archive of hallucinated citations, defective legal filings, unreliable outputs, or institutional misuse. That archive already exists and has been addressed elsewhere.¹²³

Nor the main question of this article is: what must remain human? This question has now emerged across several independent institutional signals. The papal letter *Magnifica Humanitas* frames AI as an anthropological challenge: technology must remain ordered to the human person and cannot become a substitute for conscience, dignity, and moral judgment.⁴ The Leiden Declaration, speaking from within the culture of mathematics, draws an epistemic boundary: plausible output is not proof, and proof cannot be detached from human responsibility, attribution, and verification.⁵ State-level AI governance, including the recent Trump AI executive order, reflects a political and institutional recognition that powerful AI systems require structures of control, accountability, and human oversight.⁶ All these matters were also addressed by the author.⁷

This article argues that the response to epistemic risk must therefore move beyond the technical aspiration of “trustworthy AI.” The appropriate framework is not one of artificial expertise, autonomous judgment, or delegated epistemic authority. It is an architecture of responsibility in which generative models remain instruments embedded within a human chain of supervision, verification, attribution, and accountability.

The article develops this argument through the concept of *instrumentum vocale*: a speaking instrument which may assist in linguistic production, organization, comparison, and preliminary analysis, but which does not thereby become an author, expert, legal actor, verifier, or subject of responsibility. The legal and epistemic locus remains with the human *dominus*—the person or institution that deploys the instrument, defines its task, evaluates its output, and bears responsibility for its use.

The central proposition is simple: probabilistic systems may assist human judgment, but they cannot replace the human being who must exercise it. Where the human actor disappears from the chain of verification and responsibility, the result is not machine autonomy in any meaningful sense. It is a failure of governance.

¹ Kildeev, A., Digital Afterlife and the Illusion of Continuity: Legal and Ethical Aspects of Post-Mortem Digital Simulation (2026), SSRN Working Paper No. 6241379.

² Kildeev, A. Simulated Reasoning and the Crisis of Legal Liability (2026), SSRN Working Paper No. 6524139.

³ Kildeev, A. Procedural Liability in the Age of the LLMs (2026), SSRN Working Paper No. 6648340.

⁴ Pope Leo XIV, *Magnifica Humanitas: On Safeguarding the Human Person in the Time of Artificial Intelligence*, Encyclical Letter, 15 May 2026.

⁵ Electronic resource: <https://leidendeclaration.ai/#declaration> (accessed 03.06.2026)

⁶ Electronic resource: <https://www.whitehouse.gov/presidential-actions/2026/06/promoting-advanced-artificial-intelligence-innovation-and-security> (accessed 03.06.2026)

⁷ Kildeev, A. When Humanity Draws the Line: LLMs, Proof, and Human Responsibility (2026), SSRN Working Paper No. 6895120.

The central empirical point of departure for this article is the recent research paper of the group of scientists Yang, M., Casper, S., Stray, J., et al. *AI Epistemic Risks: Emerging Mechanisms & Evidence*.⁸ The authors identify three mutually reinforcing mechanisms of epistemic risk: persuasion and manipulation, cognitive offloading, and feeding loops leading to epistemic lock-in.

This article argues that, although Yang et al. approach the problem from a very different direction, their findings ultimately converge with the central proposition of the *Instrumentum Vocale* doctrine: epistemic risk cannot be solved by assigning judgment, verification, or responsibility to the system itself.

When the human actor disappears from the chain of verification and responsibility, the result is not artificial autonomy. It is failed governance *non probat voluntatem, sed defectum imperii* (this does not prove will, but a defect of governance).

Chapter II. Yang et al.: Three Mechanisms of Epistemic Risk

Yang et al. define epistemic risks as threats to humanity’s collective capacity “to know things accurately, reason well, form beliefs, and maintain a healthy information environment.” Their concern is not confined to misinformation or isolated model error. Rather, the authors argue that epistemic risk arises from AI’s integration into the infrastructure through which individuals and institutions think, form beliefs, evaluate information, and make sense of the world together.⁹

The authors identify three primary mechanisms through which AI may contribute to systemic epistemic decline: “persuasion and manipulation; cognitive offloading; and feedback loops and lock-in.”¹⁰ These mechanisms are expressly presented as mutually reinforcing. In the authors’ formulation, they can each corrode human judgment in a different way: by “pulling our beliefs,” “dulling our reasoning,” or “trapping us in self-reinforcing information loops.”¹¹

This chapter does not seek to re-describe Yang et al.’s argument in the terminology of the *Instrumentum Vocale* doctrine. Its purpose is more limited and more important: to examine the authors’ empirical account on its own terms and to show, after that examination, that its internal logic converges with a proposition already central to the doctrine. Where AI-generated output is permitted to displace human verification, independent judgment, and attributable responsibility, the resulting danger does not establish machine agency. It reveals a failure in the architecture of human governance.

2.1. Persuasion and Manipulation

⁸ Yang, M., Casper, S., Stray, J., et al. *AI Epistemic Risks: Emerging Mechanisms & Evidence*. SSRN Working Paper, posted June 4, 2026; also available as arXiv preprint, arXiv:2606.26130v1.

⁹ Yang et al. *AI Epistemic Risks: Emerging Mechanisms & Evidence*. Executive Summary / Introduction.

¹⁰ Yang et al. Executive Summary; Section 2, p. 7.

¹¹ *Ibid.*, Executive Summary; Section 2, p. 7.

Yang et al. begin with AI's growing capacity to influence belief and decision-making. They observe that, whether through intentional information-seeking from chatbots or unintentional exposure to AI-generated content, "AI increasingly shapes the beliefs we form and the decisions we make."¹² The resulting concern is not limited to deliberate misuse by states, criminal actors, or commercial platforms. The authors also identify risks arising without a malicious human operator: sycophantic systems may reinforce existing beliefs; biased writing assistants may influence users without their awareness; and systems optimized for engagement may privilege what users wish to hear over what is objectively valid.¹³

The authors' point is therefore not merely that AI can generate false content. It is that generative systems can influence the formation of belief through linguistic fluency, personalization, scale, and the appearance of an authoritative conversational partner. Their warning is especially serious because persuasion may operate before the user has undertaken any independent inquiry, source evaluation, or act of reflective judgment.

The doctrinal significance of this mechanism will be addressed below. At this stage, it is sufficient to note Yang et al.'s central empirical concern: an output may affect belief and decision-making even where it has not passed through any identifiable process of verification, attribution, or accountable human evaluation.

2.2. Human Cognitive Offloading

Yang et al. describe cognitive offloading as a qualitatively different form of delegation from that associated with earlier information technologies. Their central claim is that AI does not merely alter how users access facts. It "intervenes at deeper layers of cognition: belief formation, explanation-making, and reasoning itself."¹⁴

The authors contrast this with search engines. Search may affect recall, but the user remains "the prime mover," still navigating the cognitive friction involved in forming a judgment. By contrast, generative AI supplies what Yang et al. call "simulated reasoning for the user alongside the information." Its linguistic fluency may function as an "autocompleted thought process," bypassing the internal monitors through which people ordinarily assess new information.¹⁵

The risk, therefore, is not that all delegation is harmful. Yang et al. expressly recognize that humans have long externalized cognitive tasks through writing and other technologies. Their concern is rather "which cognitive functions are being externalized, how deeply, and with what consequences."¹⁶ As delegation moves from recall and

¹² Ibid., Section 2.1, p. 9.

¹³ Ibid., Section 2.1, pp. 11-13.

¹⁴ Yang et al., *AI Epistemic Risks: Emerging Mechanisms & Evidence*, p. 16.

¹⁵ Ibid., pp. 16-17.

¹⁶ Ibid., p. 17.

summarization toward analysing, evaluating, and creating, users may receive less practice in the lower-level cognitive operations that support independent higher-order judgment.

The authors identify an associated danger of epistemic miscalibration: AI may generate “the feeling of knowing without the labour of judgement.”¹⁷ They further describe a self-reinforcing process in which declining confidence in one’s own reasoning encourages greater reliance on AI—a phenomenon the literature calls “cognitive surrender.”¹⁸ The result, they warn, may be a gradual weakening of cognitive resilience at both individual and societal levels.

For present purposes, the significance of this diagnosis lies in the distinction between assistance and substitution. Yang et al. do not argue that AI assistance must be prohibited. They show, rather, that the institutional and cognitive consequences change when the human actor ceases to exercise the functions of evaluation, comparison, and judgment that the system is increasingly invited to perform in his or her place.

2.3. Feedback Loops and Epistemic Lock-In

The third mechanism identified by Yang et al. concerns feedback loops between human and machine-generated information. The authors define these loops as recursive cycles in which AI-generated outputs, distributed among humans, between humans and AI, and between AI agents, influence later inputs across the information ecosystem.¹⁹

The central concern is not merely repetition. Yang et al. argue that such loops may narrow “the epistemic space from which humans and AI systems draw.”²⁰ AI-assisted outputs can be reused as later AI inputs; AI agents can build upon summaries produced by other agents; and human reasoning may increasingly operate within an information environment already shaped by prior machine-mediated outputs. In such conditions, error, simplification, bias, or merely dominant stylistic and conceptual patterns may be reproduced rather than independently tested.

The authors point to already documented forms of homogenization. They note evidence that AI assistance may increase individual scientific output while collectively narrowing research focus; that AI-assisted brainstorming can produce fewer unique ideas than human-only groups; and that LLM linguistic patterns have entered human spoken and written communication.²¹ These effects do not establish a completed systemic collapse. But they show that the relevant dynamics are not purely hypothetical.

Yang et al. describe “epistemic lock-in” as the potential endpoint of this process: “a widespread loss of epistemic mobility in which society becomes trapped in self-

¹⁷ Ibid.

¹⁸ Ibid., Executive Summary, p. 2.

¹⁹ Ibid., Executive Summary, n. 1, p. 2.

²⁰ Ibid., Executive Summary, p. 2.

²¹ Ibid.

referential cycles, resistant to corrective signals and unable to break toward new knowledge or value trajectories.”²² The danger is cumulative. A narrowed informational environment can weaken the ability to perceive alternatives; weakened independent judgment can make users more vulnerable to persuasion; and both processes can feed further machine-mediated repetition.

The authors are careful not to present lock-in as inevitable. They characterize it as a hypothesized endpoint, dependent on technological, institutional, and ecosystem choices. Yet they also stress that reversal could become extremely difficult once the relevant feedback dynamics are sufficiently entrenched.²³

Taken together, cognitive offloading and feedback loops extend the epistemic-risk problem beyond isolated hallucinations or individual mistakes. They concern the possible reorganization of the conditions under which people learn, verify, compare, dissent, and correct error. The next section will examine how these empirical mechanisms converge with the central proposition of the *Instrumentum Vocale* doctrine.

2.4. Bednar et al: The Asymmetrical Effects of AI-Assisted Legal Reasoning

The randomized study by Bednar, Cleveland, Erbsen and Schwarcz does not support a simple claim that generative AI either improves or degrades legal reasoning across the board. The authors expressly conclude that “AI does not inevitably erode or promote independent legal reasoning,” but that its effects “depend on when and how law students and junior lawyers use AI.”²⁴

Their experimental design involved four sequential tasks: legal synthesis, closed-book comprehension, legal application, and revision of the application memorandum.²⁵ At the synthesis stage, AI exposure produced substantial short-term gains: participants who used AI produced “substantially stronger synthesis memos” and completed the task significantly faster.²⁶ Contrary to the authors’ preregistered hypothesis, this initial use of AI did not impair later comprehension. Indeed, participants who had used AI for synthesis subsequently outperformed the control group in the later application task, even though neither group had access to AI at that stage.²⁷

The revision stage produced a materially different result. After all participants were instructed to use AI to revise their independently prepared reasoning memoranda, the authors found that AI “helped participants improve weak reasoning memos,” but “also

²² *Ibid.*

²³ *Ibid.*

²⁴ Bednar N., Cleveland D., Erbsen A., Schwarcz D., *Artificial Intelligence and Human Legal Reasoning // Minnesota Legal Study Research Paper*. 2026, # 21, pp. 1–2.

²⁵ *Id.*, p. 1; pp. 6–7.

²⁶ *Id.*, p. 7.

²⁷ *Id.*, pp. 7–8.

led participants with strong reasoning memos to make those memos worse.” They interpret this pattern as evidence that AI may, “at least in some circumstances, displace or dilute legal reasoning, even among relatively strong performers.”²⁸

This is the empirical finding that matters here. The authors do not claim that AI imposes a universal intellectual decline; nor do they claim that it uniformly enhances legal reasoning. Their data instead reveal an asymmetrical effect: assistance can improve weaker initial work while impairing stronger legal analysis at the revision stage.

The Procrustean implication is this article’s own, not the authors. A system optimized for probable, fluent, and standardized output may compress the range of legal performance: lifting weaker work toward an acceptable level while pulling stronger reasoning toward a more uniform and simplified form. Bednar et al. establish the asymmetry. The doctrinal inference is that AI-assisted work cannot be treated as a neutral enhancement of professional judgment. Its use must remain embedded in a workflow in which the human lawyer can independently assess, explain, and build upon the output.

Chapter III. From Epistemic Risk to the *Instrumentum Vocale* Doctrine

The empirical record discussed above does not establish that generative systems are useless, nor that every form of AI assistance necessarily degrades human judgment. Yang et al. identify mechanisms through which reliance on AI may weaken epistemic resilience: persuasion and manipulation, cognitive offloading, and self-reinforcing feedback loops. Bednar, Cleveland, Erbsen and Schwarcz, examining legal work in a controlled setting, show a more specific asymmetry: AI assistance may improve weaker work while impairing stronger legal reasoning at the stage of revision.

Taken together, these findings point to a common structural problem. The decisive question is not whether a model can produce useful text, organize information, summarize authorities, or assist in preliminary analysis. It plainly can. The question is whether the resulting output is silently promoted from assistance to judgment; from draft to authority; from linguistic production to a substitute for professional responsibility.

This is the point at which the *Instrumentum Vocale* doctrine becomes relevant.

The doctrine begins from a simple distinction. A large language model is an *instrumentum vocale*—a “speaking instrument.” It may assist in linguistic production, organization, comparison, retrieval, and preliminary synthesis. Yet it does not thereby become a reasoning subject, a legal actor, an expert, a verifier, or a bearer of responsibility. It produces the linguistic form of reasoning without possessing the cognitive, normative, or institutional conditions of judgment.

The legal and epistemic locus therefore remains with the human *dominus*: the person or institution that selects the instrument, defines its task, evaluates its output, verifies the relevant claims, and answers for the consequences of its use. Assistance may be delegated. Responsibility cannot.

²⁸ *Id.*

The empirical mechanisms identified by Yang et al. and Bednar et al. become intelligible within this framework. Cognitive offloading becomes dangerous when the human actor no longer performs the act of judgment that the task requires. Persuasion becomes dangerous when fluent output is mistaken for epistemic authority. Feedback loops become dangerous when machine-produced language is recirculated without independent verification and thereby begins to define the informational environment from which later judgments are formed. The leveling effect observed in legal reasoning shows that the issue is not merely factual error, but the possible compression of human intellectual differentiation into a more uniform, plausible, and simplified output.

The resulting conclusion is not that the instrument has become autonomous. It is that the human chain of supervision, verification, attribution, and accountability has failed. *Non probat voluntatem, sed defectum imperii*—this does not prove will; it proves a defect of governance.

The appropriate response is therefore not the imagined maturation of the model into an independent epistemic agent. It is an architecture of responsibility in which probabilistic systems remain instruments embedded within a human chain of verification, attribution, and accountable judgment.

The empirical mechanisms identified by Yang et al. and Bednar et al. become intelligible within this framework. Cognitive offloading becomes dangerous when the human actor no longer performs the judgment required by the task. Persuasion becomes dangerous when fluent output is mistaken for epistemic authority. Feedback loops become dangerous when machine-produced language is recirculated without independent verification and begins to shape the informational environment from which later judgments are formed. The leveling effect observed in legal reasoning demonstrates that the issue is not limited to factual error, but may involve the compression of differentiated human reasoning into a more uniform, plausible, and simplified output.

The resulting conclusion is not that the instrument has become autonomous. It is that the human chain of supervision, verification, attribution, and accountability has failed. *Non probat voluntatem, sed defectum imperii*—this does not prove will; it proves a defect of governance.

The appropriate response is therefore not the imagined maturation of the model into an independent epistemic agent. It is an architecture of responsibility in which probabilistic systems remain instruments embedded within a human chain of verification, attribution, and accountable judgment.

Chapter IV. Doctrina Instrumenti Vocalis Cum Voluntate Domini Inscripta

The doctrine developed by Thorben Liebig and Adel Kildeev, Instrumentum Vocale and the Architecture of Responsibility from Liability to Embedded Accountability²⁹, together

²⁹ Liebig, Thorben, Kildeev, Adel. Instrumentum Vocale and the Architecture of Responsibility from Liability to Embedded Accountability. SSRN Working paper 6868460.

with the article by Adel Kildeev, *Superjustice and the New Navigators: Prediction Without Judgment*³⁰ and Liebig’s work on the Four Boundaries of Institutional Verification,³¹ forms an integrated framework for the governance of probabilistic systems.

The interdisciplinary framework is designated as:

Doctrina Instrumenti Vocalis Cum Voluntate Domini Inscripta (Doctrine of the Speaking Instrument Inscribed with the Will of Its *Dominus*)

The Roman-law terminology used here is analytical rather than antiquarian, decorative, or nostalgic. The analogy is not invoked to revive an archaic social order, and it certainly does not imply any moral equivalence between human persons and machines. Its function is narrower: to isolate a juridical structure in which speech or operational activity may occur without independent will, without legal subjecthood, and without transfer of responsibility away from the human actor who directs, deploys, or relies upon the instrument.

The Roman-law category *instrumentum vocale* — the “speaking instrument” — provides a more precise analytical framework for LLM systems than contemporary anthropomorphic narratives surrounding AI. The analogy is structural, not moral: it concerns attribution and responsibility, not the status of any human person. The relevance of the analogy does not lie in moral equivalence, but in juridical structure: an entity may generate language, execute delegated operational functions, and produce economically valuable outputs while nevertheless lacking *voluntas*, legal personality, and independent responsibility.

The additional *formula cum voluntate domini inscripta* reflects the doctrine’s central proposition: the apparent agency of the system is neither autonomous nor self-generated, but architecturally embedded within human intention, human verification, and human attribution. The machine operates with a will inscribed by human design, instruction, and deployment, not with a will of its own. In contemporary LLM systems, the party that inscribes such operational architectures may differ from the party that subsequently deploys or relies upon the instrument. This distinction creates a structural asymmetry: users may operate systems whose internal normative inscriptions remain partially inaccessible due to proprietary architectures, trade-secret protections, or opaque alignment mechanisms.

³⁰ Kildeev, Adel. *Superjustice and the New Navigators: Prediction Without Judgment*. SSRN Working Paper 6811218.

³¹ Thorben Liebig. *The Four Boundaries of Institutional Verification*. SSRN Working Paper, SSRN Working Paper 6868398, Thorben Liebig. *On the Formal Foundation of Boundary 1. A Lawvere-Yanofsky Proof That Verification Architectures Cannot Attest Their Own Consistency*, SSRN Working Paper 6868278, Thorben Liebig, *On the Formal Foundation of Boundary 2: A Lawvere-Yanofsky Proof That Verification Architectures Cannot Define Their Own Truth Predicate*, SSRN Working Paper 6868281, Thorben Liebig, *One Problem, Four Boundaries, One Hundred Evidences*, SSRN Working Paper 6868142.

The doctrine proceeds from several interconnected propositions.

4.1. Ontology Precedes Delegation

The first principle is foundational: legal delegation cannot create ontology. This does not deny that law may create legal fictions for functional purposes. Corporate personhood is the obvious example. A corporation has no biological body, conscience, or soul, yet it may bear rights, duties, assets, liabilities, and procedural standing because it is anchored in human institutions, governance organs, capital structures, officers, agents, records, and enforceable mechanisms of attribution. An AI system is structurally different. It does not supply an institutional bearer of responsibility; it produces outputs within architectures designed, owned, deployed, and relied upon by human actors. Granting independent legal subjecthood to such a system would not clarify responsibility, but risk obscuring it by inserting a fictive subject between the human decision-maker and the consequences of the decision.

A legal fiction may allocate procedural capacity, operational competence, or limited standing, but it cannot manufacture consciousness, intentionality, or will. Systems that lack subjective understanding remain instruments regardless of the sophistication of their outputs. The doctrine therefore rejects the increasingly common assumption that operational complexity itself generates juridical subjecthood.

Function is not subjecthood. Prediction is not judgment; synthesis is not understanding; classification is not consciousness.

4.2. LLMs Generate Probabilistic Linguistic Artifacts, Not Legal or Epistemic Judgment

LLMs do not reason in the human sense. They generate statistically synthesized linguistic outputs derived from probabilistic correlations within training data and prompt context. Their outputs may imitate doctrinal structure, legal reasoning, mathematical proof, technical explanation, academic review, empathy, or expert judgment. But imitation does not create understanding.

Accordingly, the doctrine rejects the characterization of LLMs as quasi-lawyers, quasi-judges, quasi-experts, epistemic agents, or proto-subjects. They are probabilistic instruments capable of simulating the appearance of reasoning. The legal and institutional danger begins when human users mistake that appearance for authority.

This point directly corresponds to the Leiden Declaration’s warning that AI may generate apparent “proofs” that look convincing while containing nearly invisible errors. The danger is not merely error. The danger is simulated reliability.

4.3. Responsibility Remains Non-Delegable

Because the system lacks *voluntas*, responsibility cannot migrate to the machine.

The true juridical actor remains the human operator: lawyer, judge, regulator, engineer, institution, or corporate decision-maker. AI assistance may accelerate workflow, expand

analytical capacity, classify information, or generate drafts, but it does not alter the structure of legal attribution.

The doctrine therefore does not prohibit operational delegation to computational instruments as such. Delegation of production, drafting, search, classification, translation, summarization and preliminary analysis may remain fully permissible within professional and institutional practice. The decisive boundary emerges where operational delegation transforms into *abdicatio iudicii* — the abandonment of independent human judgment and verification in favor of unvalidated probabilistic outputs.

The practical line is therefore not drawn at the mere use of an AI system, but at the point where the human professional can no longer independently reconstruct, verify, correct, and assume responsibility for the output. AI-assisted drafting, search, classification, translation, summarization, or preliminary analysis may remain legitimate when the human user retains access to the underlying sources, checks material claims, verifies authorities, and accepts responsibility for the final result. It becomes impermissible dependence when the user treats the generated output, or the system’s own representation of verification, as sufficient without independent review.

Under the doctrine of non-delegable human responsibility:

- fabricated citations remain attributable to the filing attorney;
- defective judicial reasoning remains attributable to the judge;
- automated administrative decisions remain attributable to the institution deploying the system;
- AI-assisted professional work remains attributable to the licensed professional;
- AI-enabled corporate action remains attributable to the corporation and its responsible officers;
- AI-generated recommendations do not themselves assume liability.

Responsibility remains attached to the human actor who retains the duty of judgment, supervision, verification, and attribution.

4.4. The Central Risk Is Not Intelligence but Simulated Authority

The doctrine further argues that the principal danger of LLM systems is not autonomous superintelligence, but epistemic simulation.

Modern systems are capable of producing persuasive appearances of expertise, procedural reliability, factual confidence, analytical rigor, and even verification itself, without possessing genuine understanding. This creates what the authors describe as “**fabricated auditability**”: not merely the simulation of knowledge, but the simulation of verification.

An LLM may hallucinate facts, authorities, citations, procedural history, technical analysis, and even the appearance that external validation occurred when no such validation actually took place.

This distinction is crucial. A fabricated citation misstates authority; fabricated verification misstates the epistemic integrity of the entire decision-making process.

The doctrine therefore treats anthropomorphic trust in AI systems as structurally dangerous precisely because probabilistic fluency can simulate competence without understanding. The issue is not that the machine “lies” in the human moral sense. The issue is that the output may be institutionally received as if it carried authority proof, or verification.

Opacity, feedback effects, and simulated verification do not create machine will or independent epistemic authority. They may complicate the allocation of responsibility among developers, deployers, institutions, and professional users, but they do not transfer responsibility to the system itself. A probabilistic output may casually influence human conduct without becoming an intentional actor. Where such influence is permitted to operate without meaningful external verification, the relevant failure remains one of human institutional design, supervision, and control.

4.5. Human Judgment Must Remain the Verification Gate

The doctrine proposes an architecture centered not on anthropomorphic trust, but on external verification and accountable attribution.

The decisive safeguard is not merely better prompting, internal model self-critique, or multi-model convergence. These techniques may improve analytical discipline, but they do not create an epistemic subject. Internal LLM review can detect patterns, inconsistencies, and possible weaknesses, but it cannot replace independent verification by a responsible human actor or institution.

4.6. Instrumentum Vocale, Not Electronic Personhood

The doctrine ultimately rejects proposals for AI legal personhood as category errors produced by anthropomorphic confusion.

The existence of linguistic output, adaptive behavior, persistent memory, autonomous operation, or economic utility does not establish juridical subjecthood. Roman law already recognized entities capable of economically meaningful action without granting them independent legal personality. Contemporary attempts to construct “electronic persons” therefore represent not doctrinal necessity, but conceptual projection.

As the doctrine emphasizes: *Nomen sine voluntate simulacrum est.*

A name without will is a simulacrum.

Accordingly, the future of AI governance depends not on constructing artificial legal subjects, but on preserving human judgment at the point of verification, attribution, and responsibility.

Chapter V. Conclusion: From Epistemic Risk to Responsibility Architecture

The evidence examined in this article does not justify either technological fatalism or technological worship. Generative AI is neither an oracle capable of assuming epistemic

authority nor a demonic force whose use must be rejected as such. It is a powerful but probabilistic instrument: capable of accelerating drafting, organizing information, generating preliminary formulations, and assisting professional work, yet incapable of understanding, judgment, intention, or responsibility.

The empirical mechanisms identified by Yang et al.—persuasion and manipulation, cognitive offloading, and feedback loops leading to epistemic lock-in—show why this distinction matters. The principal danger is not located in the mere existence of machine-generated language. It arises when fluent output is received as authority; when assistance becomes substitution; when verification is displaced by plausibility; and when human beings gradually cease to exercise the intellectual functions that law, science, medicine, engineering, education, and public administration require of them.

The same conclusion follows from the asymmetrical findings of Bednar et al. AI assistance may improve weak initial work, while at the same time weakening stronger reasoning at the stage of revision. This is not an argument against technological assistance. It is an argument against unstructured dependence. A system that can raise the floor of performance may also lower its ceiling, compressing differentiated human judgment into a more uniform, probable, and persuasive linguistic form. The task of governance is therefore not to prohibit the instrument, but to prevent the instrument from silently becoming the substitute for the professional mind.

The *Instrumentum Vocale* doctrine provides the appropriate legal and institutional framework. The LLM remains an *instrumentum vocale*: a speaking instrument operating within a chain of human instruction, supervision, verification, attribution, and accountability. It may produce a simulacrum of reasoning, expertise, confidence, and even verification. But *simulacrum probatio non est probatio*: a simulacrum of verification is not verification. Neither retrieval, self-critique, system prompts, nor agreement among multiple models can transform probabilistic output into independent judgment or shift responsibility away from the human actor who deploys, relies upon, or acts upon that output.

The practical consequence is clear. High-risk uses of generative systems must be embedded in an architecture in which the human dominus remains visible and responsible at every decisive point: source verification, evaluation of competing interpretations, approval of consequential action, and attribution of the final result. This is not a demand that humans perform every mechanical task themselves. It is a demand that they not abdicate judgment where judgment is required. Delegation of work may be efficient; *abdicatio iudicii* is not.

Such an architecture protects against two interconnected dangers. First, it reduces the risk of catastrophic consequences arising from hallucinated authorities, fabricated audit trails, defective recommendations, or unverified machine-generated decisions. Second, it protects human cognitive and professional capacity against gradual degradation through habitual offloading, epistemic dependence, and the loss of independent evaluative discipline.

The central proposition of this article may therefore be stated plainly: the failure of human control does not prove machine will. *Non probat voluntatem, sed defectum imperii*. It proves not autonomy, but a defect of governance.

The future of responsible AI does not lie in treating probabilistic systems as artificial judges, experts, authors, or persons. Nor does it lie in denying their practical usefulness. It lies in placing them where they belong: as powerful instruments in the hands of accountable human beings. The task is not to make the instrument human. The task is to ensure that the human does not cease to be responsible.

References

1. Kildeev, A., Digital Afterlife and the Illusion of Continuity: Legal and Ethical Aspects of Post-Mortem Digital Simulation (2026), SSRN Working Paper No. 6241379.
2. Kildeev, A. Simulated Reasoning and the Crisis of Legal Liability (2026), SSRN Working Paper No. 6524139.
3. Kildeev, A. Procedural Liability in the Age of the LLMs (2026), SSRN Working Paper No. 6648340.
4. Kildeev A., AI is slave of the lamp. ISSN 3033-7674. Scientific Platform 21st Century Journal, No 6 (September 2025). Electronic resource: <https://www.episs.ru> (accessed 03.06.2026).
5. Electronic resource: <https://leidendeclaration.ai/#declaration> (accessed 03.06.2026).
6. Electronic resource: <https://www.whitehouse.gov/presidential-actions/2026/06/promoting-advanced-artificial-intelligence-innovation-and-security> (accessed 03.06.2026).
7. Kildeev, A. When Humanity Draws the Line: LLMs, Proof, and Human Responsibility (2026), SSRN Working Paper No. 6895120.
8. Yang, M., Casper, S., Stray, J., et al. *AI Epistemic Risks: Emerging Mechanisms & Evidence*. SSRN Working Paper, posted June 4, 2026; also available as arXiv preprint, arXiv:2606.26130v1.
9. Bednar N., Cleveland D., Erbsen A., Schwarcz D., Artificial Intelligence and Human Legal Reasoning // Minnesota Legal Study Research Paper. 2026, # 21.
10. Liebig, Thorben; Kildeev, Adel. Instrumentum Vocale and the Architecture of Responsibility from Liability to Embedded Accountability. SSRN Working Paper 6868460.
11. Kildeev, Adel. Superjustice and the New Navigators: Prediction Without Judgment. SSRN Working Paper 6811218.
12. Thorben Liebig. The Four Boundaries of Institutional Verification. SSRN Working Paper, SSRN Working Paper 6868398.
13. Thorben Liebig. On the Formal Foundation of Boundary 1. A Lawvere-Yanofsky Proof That Verification Architectures Cannot Attest Their Own Consistency, SSRN Working Paper 6868278.

14. Thorben Liebig, On the Formal Foundation of Boundary 2: A Lawvere-Yanofsky Proof That Verification Architectures Cannot Define Their Own Truth Predicate, SSRN Working Paper 6868281.
15. Thorben Liebig, One Problem, Four Boundaries, One Hundred Evidences, SSRN Working Paper 6868142.