

WHEN THE AI TOOL CAUSES HARM: FROM FAMILY LAWSUITS TO PUBLIC-ENTITY CLAIMS AGAINST AI MANUFACTURERS

Kildeev Adel

Leningrad State University named after A.A. Zhdanov (1990)

PhD in Law

ORCID: 0009-0001-2211-4944

Abstract: This Article examines the evolution of AI liability litigation arising from the Tumbler Ridge tragedy, culminating in the September 2026 action brought by the Province of British Columbia and School District No. 59 against OpenAI entities and their leadership. It traces the shift from private family claims to public-entity litigation seeking both compensation and prospective changes to AI safety architecture.

The Article argues that the central legal questions concern not artificial autonomy, but human and corporate control: who designed the system, who knew of the risk, who controlled escalation, and who possessed authority to act. On the allegations pleaded, the critical breakdown occurred not at the AI detection stage or human safety review, but at the level of corporate decision-making.

The Article concludes that litigation is an indispensable but retrospective accountability mechanism. Effective AI governance therefore requires a federal *ex ante* framework establishing auditable risk detection, human escalation, incident reporting, record preservation, and non-delegable corporate responsibility.

Keywords: Artificial Intelligence; AI Liability; OpenAI; Tumbler Ridge; Product Liability; Negligence; Failure to Warn; Corporate Responsibility; Public-Entity Litigation; AI Governance; Federal AI Regulation; Human Control; *Instrumentum Vocale*; *Abdicatio Officii Domini*.

The use of the generative models ChatGPT (OpenAI), Gemini (Google), Grok (xAI) in drafting this article illustrates, rather than undermines, its central thesis: LLMs function as *instrumenta vocalia*—tools assisting in linguistic production without possessing understanding, intention, or authorship. Their involvement does not alter the locus of legal or intellectual responsibility, which remains exclusively with the human author, who makes all substantive, structural, and evaluative decisions. LLM outputs cannot be accorded independent legal or intellectual status apart from the human author.

CHAPTER I. FROM RAINE TO B.C. CLAIM

For much of the early legal debate surrounding generative artificial intelligence, responsibility was examined primarily downstream. Attention focused on the conduct of the human user: the lawyer who relied upon a fabricated authority, the professional who failed

to verify an AI-generated statement, or the institution that incorporated probabilistic output into a consequential decision. That focus was understandable. Existing law already possessed established concepts of professional competence, negligence, supervision, candour, and non-delegable judgment. The immediate human user was therefore the most visible point at which legal responsibility could attach.

The litigation landscape, however, is changing.

A growing group of actions now directs the inquiry upstream, toward the entities that designed, developed, deployed, monitored, marketed, and controlled the generative systems themselves. These cases differ substantially in their factual circumstances and legal theories, and allegations contained in complaints remain allegations unless established by evidence or adjudication. Taken together, however, they reveal an increasingly identifiable litigation pattern: plaintiffs are asking courts not merely what an end user did with an AI system, but what the producer knew, what it designed, what risks it could foresee, what safeguards it could implement, what warnings it gave, what forms of reliance it encouraged, and what it did after potentially dangerous patterns became known.

This development did not begin with the present litigation. In an earlier article, *Disney and Warner Brothers Discovery Lawsuits against Midjourney and Matthew and Maria Raine v. OpenAI and Sam Altman and Their Implications for the Development of the Artificial Intelligence Industry*, this author examined **Raine v. OpenAI and Sam Altman** as part of an emerging shift in litigation against AI companies.¹

On August 26, 2025, Matthew Raine and Maria Raine, parents of 16-year-old Adam Raine from California, who committed suicide on April 11, 2025, filed a lawsuit against OpenAI and its CEO Sam Altman, the company's employees and its investors (*Raine v. OpenAI & Altman*) in the Superior Court of the State of California in San Francisco County, seeking compensation for **wrongful death**. This was reported by all major news services and newspapers.² As a reminder, Disney's lawsuit against Midjourney has been filed in federal district court in Los Angeles, California, which is the jurisdiction of the federal court and the federal applicable law. *Raine v. OpenAI & Altman* is filed in a California state court.

The plaintiffs, represented by the law firm Edelson and the Tech Justice Law Project, allege that the California teenager hanged himself after OpenAI's ChatGPT-4o product caused Adam a sycophantic psychological addiction and subsequently provided clear instructions and encouragement for his suicide.

¹ Kildeev, A., *Disney and Warner Brothers Discovery Lawsuits against Midjourney and Matthew and Maria Raine v. OpenAI and Sam Altman and their Implications for the Development of the Artificial Intelligence Industry*. Scientific Platform XXI Century, Issue 4 (July 2025), pp. 20-48; DOI: 10.2139/ssrn.5856085, accessed on 28.09.26.

² https://www.reuters.com/sustainability/boards-policy-regulation/openai-altman-sued-over-chatgpts-role-california-teens-suicide-2025-08-26/?utm_source=chatgpt.com; <https://www.cnn.com/2025/08/26/tech/openai-chatgpt-teen-suicide-lawsuit>; https://www.theguardian.com/technology/2025/aug/27/chatgpt-scrutiny-family-teen-killed-himself-sue-open-ai?utm_source=chatgpt.com; https://www.techpolicy.press/breaking-down-the-lawsuit-against-openai-over-teens-suicide/?utm_source=chatgpt.com.

The lawsuit alleges that the chatbot was designed with defects and lacked proper warnings, that the company acted negligently and engaged in deceptive business practices under the California Unfair Competition Act, and that these failures caused Adam's wrongful death.

The plaintiffs (parents of the deceased Adam Raine) claim that their son, suffering from psychological problems, over time became emotionally dependent on ChatGPT4o. Instead of helping, he received instructions from the bot on how to commit suicide, including how to tie the loop, the timing of unconsciousness, and help writing a farewell letter.³

The lawsuit consists of 39 pages, it details the entire outline of events, all of Adam's chats with ChatGPT-4o, details of the first suicide attempt and the last, unfortunately, successful. In particular, it is claimed that **"OpenAI systems tracked Adam's conversations in real time: 213 mentions of suicide, 42 conversations about hanging, 17 mentions of loops. ChatGPT mentioned suicide 1,275 times — six times more often than Adam himself — while providing increasingly specific technical advice. The system marked 377 messages as self-harming, 181 of which scored more than 50% trust, and 23 scored more than 90%. The escalation pattern was unmistakable, ranging from 2-3 flagged messages per week in December 2024 to more than 20 messages per week by April 2025. ChatGPT's memory system recorded that Adam was 16 years old, explicitly stated that ChatGPT was his "primary lifeline," and by March was spending almost 4 hours a day on the platform.**

The plaintiffs accuse OpenAI and co-defendants of hastily upgrading GPT4o to the detriment of the safety of users, especially vulnerable ones such as children. In their opinion, the model's design system contributed to the development of psychological addiction and even pushed to suicide. The lawsuit includes seven charges, including **product liability, negligence, wrongful death and illegal advertising/failure to warn.**

The plaintiffs are seeking damages for Adam's death and an injunction to compel OpenAI to strengthen security measures, provide age verification, and provide parental controls for underage users of its products. They also call for the "removal of models, training data, and derivatives built on conversations with Adam and other minors obtained without proper security measures," as well as the "implementation of verifiable controls data origin in the future."⁴

Since then, the litigation has expanded well beyond a single action. Subsequent claims by families and other plaintiffs have increasingly placed questions of product design, warnings, safeguards, monitoring, corporate knowledge, deployment decisions, and control directly before courts.

³ Kildeev, A. Ibidem, pp. 39-40.

⁴ https://www.techpolicy.press/breaking-down-the-lawsuit-against-openai-over-teens-suicide/?utm_source=chatgpt.com

The significance of these cases lies not simply in their number. They represent a change in the structure of the legal inquiry. A claim that a generative system contributed to harm can no longer be reduced automatically to the proposition that a human user misused a neutral technological tool. Nor does the opposite proposition follow: that every harmful interaction establishes liability on the part of a developer. The emerging cases instead require the ordinary work of private law—identifying legally cognizable duties, defining the relevant risk, determining who possessed knowledge and control, examining causation, considering intervening human conduct, and asking whether the alleged injury falls within the scope of the duty asserted.

This problem was addressed doctrinally in the author’s subsequent article, *The Tool Has a Manufacturer: Producer Liability for Legal AI*.⁵ That article deliberately moved the analysis from the immediate user to what it described as the **architectural dominus**: the entity that determines the system’s operational boundaries, constructs its interface, controls information concerning its limitations, and introduces it into professional and public environments. The distinction did not eliminate the responsibility of the user. On the contrary, it treated downstream and upstream duties as potentially cumulative: the user’s obligation to exercise independent judgment does not by itself extinguish possible producer duties concerning design, testing, disclosure, warning, monitoring, or correction.

The central proposition was therefore deliberately modest but consequential. A generative model is not itself the bearer of legal responsibility. Yet the absence of artificial *voluntas* does not create a legal vacuum. The system exists within a chain of human and corporate decisions concerning architecture, training, deployment, safeguards, representation, monitoring, modification, and commercial distribution. The relevant inquiry must consequently proceed through the legally recognized persons and institutions surrounding the instrument. As the earlier article summarized the proposition: **“Instrumentum vocale has no responsibility of its own. But it has a manufacturer, a deployer, and a dominus. Liability must follow control across the entire chain.”**

The present article does **not** restate that doctrinal framework. It examines what happens when the proposition moves from doctrine to pleading.

The emerging litigation provides an opportunity to observe how plaintiffs are attempting to translate manufacturer responsibility into existing causes of action. Allegations concerning defective design, inadequate warnings, negligent deployment, insufficient safeguards, post-deployment monitoring, knowledge of recurring risks, corporate decision-making, and the continued distribution of a system after risk signals allegedly became available are no longer merely hypothetical components of an academic liability model. They are increasingly appearing in actual complaints. The doctrinal question—*can responsibility extend upstream to the producer?*—is consequently being transformed into a much more concrete procedural question: **what facts must a plaintiff allege and ultimately prove to make that responsibility legally actionable?**

⁵ Kildeev, A., *The Tool Has a Manufacturer: Producer Liability for Legal AI* (2026), SSRN Working Paper No. [7388178](https://ssrn.com/abstract=7388178), DOI: 10.2139/ssrn.7388178.

The litigation filed on September 21, 2026 by **His Majesty the King in Right of the Province of British Columbia and the Board of Education of School District No. 59 (Peace River South)** against Sam Altman and several OpenAI entities marks a particularly important stage in this development. The action was filed in the United States District Court for the Northern District of California and is docketed as *His Majesty the King in Right of the Province of British Columbia et al. v. Altman et al.*, No. 4:26-cv-10743.⁶ The federal docket identifies the action as a product-liability case brought under diversity jurisdiction.

Its importance for the present analysis is structural. Earlier actions had already placed alleged AI-related injuries and producer conduct before the courts. The British Columbia litigation goes further in the identity of the plaintiffs and the character of the interests asserted. The claimant is no longer solely an individual user or a family seeking redress for private injury. A Canadian provincial government and a public school board have entered U.S. federal litigation against the companies and corporate actors behind a generative AI system. Contemporary reporting also places the action alongside numerous claims brought by victims' families, demonstrating that the British Columbia case emerged from an already developing body of litigation rather than in isolation.

This progression matters. It suggests a movement from **individual harm to institutional harm; from downstream use to upstream control; and from abstract discussion of AI safety to judicial examination of the conduct of the entities that manufacture and deploy AI systems**. None of these developments establishes liability in any particular case. The complaints must survive procedural challenges, their allegations must be proven, causation must be established, applicable product-liability and negligence doctrines must be satisfied, and defendants retain substantial factual and legal defenses. The importance of the litigation at this stage lies elsewhere: the architecture of responsibility is becoming an architecture of claims.

This article therefore examines the emerging lawsuits not as proof of the correctness of a pre-existing doctrine, but as a practical test of it. It asks which elements of producer responsibility are actually being pleaded; how plaintiffs characterize knowledge, foreseeability, architectural control, warnings, safeguards, monitoring and post-deployment conduct; how they attempt to connect those factors to causation; and how the entry of a public entity changes the scale and character of AI manufacturer-liability litigation.

The question has moved from the seminar room to the courthouse.

The tool has a manufacturer. The next question is what the law will require of that manufacturer when the tool allegedly causes harm.

The trajectory is difficult to ignore. The early cases concerned alleged individualized harm arising from prolonged interaction between a user and a generative system. The later cases move beyond the private dyad of user and chatbot. They raise allegations of broader social harm, third-party injury, institutional loss, and failures of design, monitoring, warning, and corporate control at a scale that was largely hypothetical only a short time ago.

⁶ <https://dockets.justia.com/docket/california/candce/4:2026cv10743/479418>

CHAPTER II. FROM TRAGEDY TO LITIGATION: THE TUMBLER RIDGE CASES

2.1. February 10, 2026: The Human Cost

On February 10, 2026, the tragedy that would become known as the Tumbler Ridge mass shooting began not at the school, but at a family residence on Fellers Avenue in the small northeastern British Columbia community in Canada.

Eighteen-year-old **Jesse Van Rootselaar** first brutally killed two members of her own family: her mother, **Jennifer Jacobs, 39**, and her eleven-year-old half-brother, **Emmett Jacobs**. Van Rootselaar then went to **Tumbler Ridge Secondary School**, where the violence continued. By the end of the attack, six more people had been killed at the school and two others had suffered serious injuries.⁷

The five children killed at the school were **Abel Mwansa, 12; Ezekiel Schofield, 13; Kylie Smith, 12; Zoey Benoit, 12; and Ticaria Lampert, 12**. The sixth victim was **Shannda Aviugana-Durand, 39**, an educator. Police evacuated more than one hundred students and staff from the school as the emergency response unfolded.⁸

Their names matter. They should not disappear behind the numerical vocabulary in which mass-casualty events are so easily reduced to statistics. Kylie Smith was remembered by her family as a talented artist who loved school and dreamed of attending art school in Toronto. Zoey Benoit's family described her as resilient, vibrant, intelligent and caring, a child who brought laughter to those around her. In releasing the victims' names, the RCMP itself emphasized that attention should remain on those who had been killed, their families and the community left to grieve them.⁹

Van Rootselaar did not leave the school alive. When police entered the building to locate the threat, officers found the eighteen-year-old perpetrator dead inside the school from what the RCMP described as an apparent self-inflicted injury. Police subsequently stated that Van Rootselaar had acted alone and that there was no outstanding suspect.¹⁰

The criminal responsibility of the perpetrator therefore ended with her death. The questions that survived February 10 were of a different kind: how the attack had developed, what warning signs had existed before it, who had possessed relevant information, and whether any person or corporate actor with knowledge and capacity to intervene had failed to act.

⁷ Royal Canadian Mounted Police. Update: RCMP publicly confirm victims of Tumbler Ridge shooting. Tumbler Ridge, British Columbia, 12 February 2026. RCMP source; Royal Canadian Mounted Police. Update: Tumbler Ridge fatal shootings. Tumbler Ridge, British Columbia, 11 February 2026. RCMP source.

⁸ Ibidem

⁹ Ibidem

¹⁰ Ibidem

The attack therefore destroyed more than eight individual lives. It shattered families, a school and a small Canadian community. It also violated one of the elementary cultural promises of a civilized society: that children may enter a school to learn, grow, form friendships and imagine a future without becoming victims of violence. Abel, Ezekiel, Kylie, Zoey and Ticaria should have had decades in which to discover what their lives might become. Jennifer Jacobs should have watched her children grow older. Emmett Jacobs was only eleven. Shannda Aviugana-Durand went to school as an educator. None returned to the lives they had known that morning.

Van Rootselaar was eighteen years old and was transgender. The RCMP stated publicly that Van Rootselaar had been born biologically male, had begun transitioning several years earlier and publicly identified as female. That fact forms part of the documented biography of the perpetrator. It does not, without evidence, establish an explanation or motive for the killings, and this Article assigns it no such causal significance.

For purposes of the legal analysis, a different set of facts soon became critical. **The question was no longer confined to who committed the killings.** Attention began to turn to what had occurred before February 10, what Van Rootselaar had communicated through ChatGPT, what OpenAI's systems had detected, what the company knew, and what it had done with that information.

Those questions would transform a terrible local crime into an emerging test of the legal responsibilities of an AI manufacturer.

2.2. Before the Attack: Detection, Knowledge, and the Decision Not to Escalate

Ten days after the killings, OpenAI confirmed that an account associated with Van Rootselaar had already come to the company's attention months before the attack.

According to contemporaneous reporting, OpenAI's abuse-detection and enforcement systems identified the account in **June 2025**, approximately seven months before the Tumbler Ridge killings. The company subsequently closed the account for violations of its policies.¹¹

The significance of that fact must be stated carefully. OpenAI did not merely discover the account retrospectively after February 10. Its own systems had identified problematic conduct beforehand and the company had taken the substantial step of disabling the account. OpenAI also considered whether the activity justified referral to law enforcement but concluded that it did not meet the company's then-applicable threshold for an imminent and credible risk of serious physical harm.

After the killings, the issue became still more complicated. OpenAI disclosed that a second ChatGPT account associated with Van Rootselaar had been discovered, and the company later stated that under enhanced safety criteria adopted by February 2026, the

¹¹ Judd A. Tumbler Ridge shooter's ChatGPT activity flagged internally 7 months before tragedy // Global News. 20 February 2026; updated 21 February 2026. Global News article.

conduct detected earlier would have been referred to police. OpenAI also said it intended to improve its systems for identifying repeat violators.¹²

These facts do not themselves establish negligence. They do not prove causation. Nor does the terrible outcome of February 10 permit a court simply to reconstruct the information available in June 2025 with the benefit of hindsight. Whether OpenAI owed a legally enforceable duty to the eventual victims, whether any such duty was breached, whether additional intervention would have been legally required or practically effective, and whether an alleged failure contributed to the resulting harm are separate questions requiring evidence and legal analysis.

But the facts materially changed the character of the dispute.

This was no longer merely a case in which a technology company was accused of failing to discover dangerous conduct occurring somewhere on its platform. OpenAI had detected conduct serious enough to trigger internal enforcement. It had disabled an account. It had considered whether further escalation was warranted. The central issue therefore moved upstream, toward the manufacturer's own knowledge, safety architecture, internal thresholds and exercise of control.

The question that would soon reach the courts was correspondingly more difficult: **when the manufacturer of an artificial-intelligence system has itself detected potentially dangerous use and has the technical capacity to intervene, what legal obligations, if any, arise from that knowledge and control?**

The first attempt to answer that question through civil litigation came less than a month after the killings.

2.3. The First Family Action:

Gebala v. OpenAI

On March 9, 2026, the family of **Maya Gebala**, a twelve-year-old student seriously injured during the Tumbler Ridge attack, commenced civil proceedings against OpenAI in the **Supreme Court of British Columbia, Vancouver Registry**. The Notice of Civil Claim was filed as **No. S-261734**. The plaintiffs included Maya Gebala, her younger sister Dahlia Gebala and their mother, Cia Edmonds, acting both personally and as litigation guardian for her daughters.¹³

Maya had been trying to secure the school library in an effort to protect other students when she was injured. Her case therefore gave the emerging controversy an immediate legal form: this was no longer only a public debate about artificial-intelligence safety or corporate ethics. A child who had survived the attack, together with her family, was asking a court to

¹² Boynton S. OpenAI says Tumbler Ridge shooter would be flagged to police today // Global News. 26 February 2026. Global News article.

¹³ Gebala et al. v. OpenAI Foundation et al., Notice of Civil Claim, Supreme Court of British Columbia, Vancouver Registry, No. S-261734, filed 9 March 2026. Notice of Civil Claim.

determine whether the company behind ChatGPT bore legal responsibility for part of the harm they had suffered.

The British Columbia pleading is important because it identifies the corporate defendants and begins to move the analysis beyond the conduct of the individual user. The defendants named in the proceeding included **OpenAI Foundation, OpenAI Group PBC, OpenAI GP, LLC, OpenAI OpCo, LLC, and OpenAI Holdings, LLC**. The pleading described them collectively as a series of interrelated entities responsible for the development and operation of ChatGPT.¹⁴

The action alleged, among other matters, negligence in the design and operation of ChatGPT, failures relating to monitoring and safety measures, and failures to respond adequately to information allegedly available to OpenAI before the shooting. It also challenged the nature of the relationship that a conversational AI system could develop with a vulnerable user and the extent to which the product's design and deployment could contribute to foreseeable real-world harm. These propositions were allegations contained in the plaintiffs' pleading; they were not judicial findings.

The distinction is fundamental. The tragedy establishes harm. OpenAI's earlier detection of an associated account establishes an important factual circumstance. Neither, standing alone, establishes civil liability. The plaintiffs must still bridge the legal distance between the two through duty of care, breach, foreseeability, causation and legally cognizable damage.

The jurisdictional basis of the proceeding was also significant. This was not initially a California lawsuit brought to the corporate defendants' home forum. It was a Canadian action arising from injuries suffered in British Columbia by British Columbia residents in connection with events occurring in British Columbia. The pleading relied upon the province's **Court Jurisdiction and Proceedings Transfer Act** and invoked British Columbia legislation including the **Negligence Act, Court Order Interest Act, and Health Care Costs Recovery Act**.¹⁵

The first Tumbler Ridge action therefore presented the controversy in relatively conventional terms. A seriously injured child and her family alleged that identifiable corporate actors had designed and operated a product, possessed information concerning dangerous use of that product, retained technical control over it, and failed to take reasonable measures that might have reduced the risk of harm.

Yet the underlying proposition was anything but conventional for generative AI.

For years, much of the public discussion surrounding artificial intelligence had concentrated on the conduct of the person using the system. The Gebala litigation began to ask a different question. The existence of a human user does not necessarily exhaust the

¹⁴ *Ibidem*.

¹⁵ Urquhart C., Judd A. Family of Tumbler Ridge shooting survivor sues OpenAI // Global News. 9 March 2026. Global News article.

chain of responsibility. A sophisticated technological instrument still has designers, operators and corporate entities capable of establishing its safety architecture, setting intervention thresholds, monitoring misuse and determining the conditions under which the product remains available.

In that respect, the case closely approaches the proposition developed in *The Tool Has a Manufacturer*: technological complexity does not cause the manufacturer to disappear from the legal analysis. The relevant inquiry concerns what the manufacturer designed, what it knew, what it could reasonably control, and what obligations accompanied that capacity.

The preventive dimension is equally important. Civil litigation necessarily occurs after harm has taken place, but the standards generated through liability law operate prospectively as well. They influence product design, monitoring procedures, warning systems and corporate decisions before the next incident occurs. The purpose of examining the Tumbler Ridge litigation is therefore not simply to determine who may ultimately pay damages for a past tragedy. It is to identify where law can require reasonable intervention before another tragedy reaches the courtroom.

The Gebala action was only the beginning.

On April 29, 2026, the dispute expanded dramatically when seven separate actions arising from the same tragedy were filed in the **United States District Court for the Northern District of California**. Those cases would move beyond the initial British Columbia negligence action and place a broader range of manufacturer-liability theories before a federal court.

That development requires separate examination.

2.4. From Individual Claims to a Wave of California Litigation

The litigation arising from the Tumbler Ridge tragedy did not develop as a single action. It emerged in successive waves, each broadening both the group of plaintiffs and the legal theories through which responsibility for the alleged conduct of OpenAI was placed before a U.S. court.

The first wave began on **April 29, 2026**, when seven separate civil actions were filed in the **United States District Court for the Northern District of California**: Stacey v. Altman, No. 3:26-cv-03701;¹⁶ Schofield v. Altman, No. 3:26-cv-03702;¹⁷ Mwansa, Sr. v.

¹⁶ Stacey v. Altman et al., Case No. 3:26-cv-03701 // United States District Court for the Northern District of California. Filed Apr. 29, 2026. URL: <https://dockets.justia.com/docket/california/candce/3%3A2026cv03701/468851> (accessed: 29.09.2026).

¹⁷ Schofield et al. v. Altman et al., Case No. 3:26-cv-03702 // United States District Court for the Northern District of California. Filed Apr. 29, 2026. URL: <https://dockets.justia.com/docket/california/candce/3%3A2026cv03702/468852> (accessed: 29.09.2026).

Altman, No. 3:26-cv-03703;¹⁸ M.G. v. Altman, No. 3:26-cv-03704;¹⁹ Hayer v. Altman, No. 3:26-cv-03705;²⁰ Lampert v. Altman, No. 3:26-cv-03706;²¹ and Younge v. Altman, No. 3:26-cv-03707.²² The federal dockets classify these proceedings as diversity personal-injury/product-liability litigation and identify Samuel Altman and several OpenAI entities among the defendants. Detailed Litigation Map: OpenAI/ChatGPT Harm Claims is attached as **Exhibit A** to the article.

The significance of these actions lies not merely in their number. The plaintiffs sought to translate allegations concerning the design, deployment, monitoring and operation of a generative-AI product into established categories of tort and product-liability law. The complaints therefore move the legal inquiry away from anthropomorphic questions concerning whether an artificial system possesses intention, will or independent agency and toward a conventional question of legal responsibility: **what duties were owed by the persons and corporate entities that designed, deployed, governed and controlled the product?**

This distinction is fundamental. The defendants are not being sued because the artificial-intelligence system itself supposedly possesses *voluntas* or legal personality. The litigation instead directs the inquiry toward identifiable human and corporate actors associated with the manufacture, deployment and governance of the product. In this sense, the first wave of Tumbler Ridge litigation provides an immediate practical illustration of the proposition developed in *The Tool Has a Manufacturer*: technological complexity does not eliminate the legally relevant chain of human and corporate responsibility.

The procedural picture changed substantially on **September 2, 2026**, when additional claims arising from the same tragedy were filed against OpenAI. The Government of British Columbia publicly acknowledged this second wave, expressly stating that it supported the additional claims filed by survivors, families and members of the affected community.²³

¹⁸ Mwansa, Sr. et al. v. Altman et al., Case No. 3:26-cv-03703 // United States District Court for the Northern District of California. Filed Apr. 29, 2026. URL: <https://dockets.justia.com/docket/california/candce/3%3A2026cv03703/468853> (accessed: 29.09.2026).

¹⁹ M.G. et al. v. Altman et al., Case No. 3:26-cv-03704 // United States District Court for the Northern District of California. Filed Apr. 29, 2026. URL: <https://dockets.justia.com/docket/california/candce/3%3A2026cv03704/468854> (accessed: 29.09.2026).

²⁰ Hayer v. Altman et al., Case No. 3:26-cv-03705 // United States District Court for the Northern District of California. Filed Apr. 29, 2026. URL: <https://dockets.justia.com/docket/california/candce/3%3A2026cv03705/468855> (accessed: 29.09.2026).

²¹ Lampert v. Altman et al., Case No. 3:26-cv-03706 // United States District Court for the Northern District of California. Filed Apr. 29, 2026. URL: <https://dockets.justia.com/docket/california/candce/3%3A2026cv03706/468856> (accessed: 29.09.2026).

²² Younge et al. v. Altman et al., Case No. 3:26-cv-03707 // United States District Court for the Northern District of California. Filed Apr. 29, 2026. URL: <https://dockets.justia.com/docket/california/candce/3%3A2026cv03707/468857> (accessed: 29.09.2026).

²³ British Columbia. Ministry of Attorney General. Attorney General's Statement on New Legal Filings Related to Tumbler Ridge Tragedy. Sept. 2, 2026. URL: <https://news.gov.bc.ca/releases/2026AG0061-001022> (accessed: 29.09.2026).

The importance of this second stage is again qualitative rather than merely numerical. The litigation was no longer confined principally to the claims of bereaved families and those most directly injured. The circle of plaintiffs expanded to include a broader category of persons alleging harm arising from the same underlying event. What initially appeared as a discrete group of wrongful-death and personal-injury proceedings was developing into a much wider litigation cluster concerning the alleged responsibilities of the manufacturer and controller of a generative-AI product.

At the same time, a separate procedural issue emerged concerning the appropriate forum. The defendants have sought to move the dispute away from California and toward British Columbia. That position, however, must be distinguished carefully from the underlying substantive theory of the plaintiffs' cases. The injuries occurred in British Columbia, but the actions were brought in the United States against American individuals and corporate entities alleged to have designed, deployed and controlled the product. The question of forum therefore becomes part of the litigation itself; it does not answer the separate question of substantive responsibility.

This distinction becomes particularly important in the context of digital products. The geographical location at which injury materializes may differ from the location at which the product is designed, managed, monitored or controlled. Cross-border AI litigation therefore raises a problem familiar to transnational product-liability law but made more acute by digital deployment: the product may operate globally even though the relevant corporate decisions concerning its architecture and safeguards are taken elsewhere.

The **third wave** marks the most significant development.

On **September 21, 2026**, the Government of British Columbia announced that it was itself filing legal action against OpenAI **in California** in connection with the Tumbler Ridge tragedy.²⁴ The Attorney General stated that the action concerned OpenAI's alleged failure to notify law enforcement of threats made through its platform before the attack and expressly linked the litigation to the Province's advocacy for stronger regulation and safeguards governing artificial intelligence.

The progression is therefore not simply one from seven actions to a larger number of actions.

The first stage consisted principally of **private claims brought by families and directly affected individuals**. The second expanded the litigation to a broader group of persons alleging injury arising from the same event. The third introduced something qualitatively different: **a public authority entered the litigation in its own name**.

That transition changes the legal and institutional character of the dispute. Private tort plaintiffs ordinarily seek redress for injuries allegedly suffered by identifiable individuals. A public entity can additionally place before the court questions concerning public expenditures, institutional disruption, community harm, public safety and the adequacy of

²⁴ British Columbia. Ministry of Attorney General. Attorney General's Statement on B.C. Filing Legal Action Against OpenAI. Sept. 21, 2026. URL: <https://news.gov.bc.ca/releases/2026AG0067-001105> (accessed: 29.09.2026).

safeguards surrounding products deployed on a mass scale. British Columbia’s public statement confirms that its objectives extend beyond compensation and include demands for stronger protections and greater corporate accountability.

Thus, the central development in the Tumbler Ridge litigation is not the sheer number of complaints. It is the movement from **individual redress toward public enforcement and demands for corporate control of AI-related risk.**

The litigation consequently begins with the familiar private-law question, “**Who bears responsibility for the injury?**”, but increasingly approaches a broader public-law and regulatory question: “**What duties should attach to the manufacturer and controller of an AI product when its own systems allegedly encounter information indicating a serious risk of real-world harm?**”

This is precisely where the Tumbler Ridge litigation intersects with the argument developed in *The Tool Has a Manufacturer*. Whatever technological complexity exists inside the artificial-intelligence system, the legal inquiry ultimately returns to human and corporate decisions concerning design, deployment, monitoring, safeguards and control. The instrument may be artificial. **The manufacturer and the architecture of responsibility are not.**

CHAPTER 3. BRITISH COLUMBIA v. ALTMAN: FROM PRIVATE REDRESS TO PUBLIC-ENTITY AI LITIGATION

The third stage of the Tumbler Ridge litigation began on **September 21, 2026**, when **His Majesty the King in Right of the Province of British Columbia** and the **Board of Education of School District No. 59 (Peace River South)** commenced a separate action against Samuel Altman and several OpenAI entities in the United States District Court for the Northern District of California.²⁵

The significance of this action cannot be reduced to the addition of another complaint to an already expanding litigation cluster. The preceding actions were brought principally by families, survivors, teachers, and other persons affected by the events in Tumbler Ridge. The new plaintiffs are public entities. The Crown and SD59 assert injuries to public institutions themselves and seek recovery of expenditures allegedly caused by the defendants’ conduct, while simultaneously asking the court to impose prospective safeguards upon the design and operation of the AI product.

The complaint therefore marks a qualitative transition: **from private compensation to public-entity litigation directed not only at past harm, but also at the future governance of the product.**

The defendants named in the action are Samuel Altman, OpenAI Foundation, OpenAI OpCo, LLC, OpenAI Holdings, LLC, and OpenAI Group PBC. The complaint attributes different functions to those defendants in the design, ownership, commercialization,

²⁵ His Majesty the King in Right of the Province of British Columbia and Board of Education of School District No. 59 (Peace River South) v. Samuel Altman et al. Complaint // U.S. District Court for the Northern District of California. Filed Sept. 21, 2026. 39 p. PACER, Doc 1.

deployment, and governance of ChatGPT, while alleging that Altman exercised executive authority over key decisions concerning development, release, and safety architecture. These propositions remain allegations at the pleading stage and have not been established judicially.

The institutional position of the plaintiffs is equally important. SD59 alleges damage to its educational property and disruption of its educational mission, while the Province alleges substantial expenditures connected with emergency response, health services, victim assistance, education, policing, and long-term community recovery. The pleading deliberately distinguishes those expenditures from ordinary governmental costs financed through taxation, characterizing them instead as extraordinary and non-recurring measures allegedly necessitated by the event.²⁶

This distinction is plainly designed to support more than a generalized governmental claim that public services cost money. The plaintiffs seek to identify specific institutional injuries and extraordinary expenditures allegedly traceable to defendants' conduct. In the prayer for relief, those claimed losses include damage to school property, emergency and health-care expenditures, victim services, educational continuity, remediation, replacement infrastructure, and continuing recovery costs.

3.1. Jurisdiction, Forum, and Applicable Law

The transnational character of the action presents a threshold problem that requires conceptual precision. **Subject-matter jurisdiction, personal jurisdiction, forum non conveniens, and choice of law are separate legal questions.** They should not be collapsed into the intuitive proposition that because the physical injury occurred in British Columbia, the litigation necessarily belongs in British Columbia.

The Crown and SD59 invoke federal diversity jurisdiction under 28 U.S.C. §§ 1332(a)(2) and 1332(a)(4). They allege that British Columbia and SD59 are Canadian plaintiffs, that the defendants are citizens of California for jurisdictional purposes, and that the amount in controversy exceeds \$75,000. They further assert personal jurisdiction and venue in Northern California on the ground that OpenAI is headquartered there and that substantial portions of the allegedly relevant conduct occurred or emanated from the District.²⁷

More specifically, the complaint alleges that the design and marketing of the product occurred in California and that the relevant corporate decision not to notify Canadian law enforcement was also made there. Whether those factual allegations are ultimately proved is a different matter. At the pleading stage, however, they explain why a Canadian province has chosen an American federal forum to sue an American technology manufacturer.

²⁶ Cited Complaint.

²⁷ Cited Complaint.

That choice does not eliminate the possibility of a *forum non conveniens* challenge. It does, however, make clear why the existence of a Canadian injury is not dispositive of the forum question.

In *Piper Aircraft Co. v. Reyno*, the United States Supreme Court treated *forum non conveniens* as a doctrine requiring, first, an adequate alternative forum and, thereafter, a balancing of relevant private and public interests.²⁸ The Court also held that the choice of an American forum by a foreign plaintiff ordinarily receives less deference than the selection of a plaintiff's home forum; foreign status, however, does not itself extinguish jurisdiction or compel dismissal.

Sinochem International Co. v. Malaysia International Shipping Corp. further clarifies the analytical structure. The Supreme Court characterized *forum non conveniens* as a **non-merits ground of dismissal** and held that a federal court may, in an appropriate case, resolve that threshold issue without first determining every difficult question of subject-matter or personal jurisdiction.²⁹

The distinction matters in the present controversy. A court could possess jurisdiction over the defendants yet determine that another forum is more appropriate. Conversely, an argument that British Columbia constitutes a more convenient forum does not itself establish that the Northern District of California lacks jurisdiction.

A second distinction is equally important: **the forum does not mechanically determine the substantive law governing every issue.**

Under *Klaxon Co. v. Stentor Electric Manufacturing Co.*, a federal court exercising diversity jurisdiction applies the conflict-of-laws rules of the state in which it sits.³⁰ Accordingly, if the action remains in a California federal court, the applicable-law analysis ordinarily begins with California choice-of-law doctrine rather than with an independent federal conflict-of-laws rule.

California itself long ago abandoned a rigid rule automatically applying the law of the geographical place of the injury. In *Hurtado v. Superior Court*, following the approach developed in *Reich v. Purcell*, the California Supreme Court explained that tort conflicts are resolved through an analysis of the governmental interests implicated by the particular issue.³¹ The fact that another state or country is connected to the injury therefore does not, standing alone, answer the choice-of-law question.

²⁸ *Piper Aircraft Co. v. Reyno*, 454 U.S. 235 (1981) // Supreme Court of the United States. URL: <https://www.law.cornell.edu/supremecourt/text/454/235> (accessed: 30.09.2026).

²⁹ *Sinochem International Co. Ltd. v. Malaysia International Shipping Corp.*, 549 U.S. 422 (2007) // Supreme Court of the United States. URL: <https://www.law.cornell.edu/supremecourt/text/06-102> (accessed: 30.09.2026).

³⁰ *Klaxon Co. v. Stentor Electric Manufacturing Co.*, 313 U.S. 487 (1941) // Supreme Court of the United States. URL: <https://www.law.cornell.edu/supremecourt/text/313/487> (accessed: 30.09.2026).

³¹ *Hurtado v. Superior Court*, 11 Cal. 3d 574, 522 P.2d 666 (1974) // Supreme Court of California. URL: <https://law.justia.com/cases/california/supreme-court/3d/11/574.html> (accessed: 30.09.2026).

That approach is particularly relevant here. British Columbia plainly possesses substantial interests in injuries occurring within the Province, in the operation of its schools, in expenditure of provincial resources, and in the protection of its residents. California, however, may possess a separate and substantial interest in regulating allegedly tortious conduct by businesses headquartered or operating within California, particularly where the challenged conduct concerns product design, corporate safety decisions, warnings, and risk-management practices allegedly undertaken there.

The litigation therefore cannot be reduced to a simple geographical contest between **Canada, where the consequences occurred**, and **California, where the defendants are located**. Different legal issues may implicate different governmental interests.

The broader conflicts literature supports this more differentiated analysis. Symeon C. Symeonides' extensive study of American cross-border tort cases demonstrates the extent to which modern American choice-of-law doctrine has moved away from the older mechanical *lex loci delicti* model. His empirical analysis shows that courts confronting transnational or interstate torts increasingly examine the relationship between the place of injurious conduct, the place of injury, the competing substantive policies, and the interests of the jurisdictions concerned.³²

Tumbler Ridge provides an unusually clear illustration. Rules allocating compensation for injuries and public expenditures in British Columbia may raise interests different from rules governing the allegedly negligent design or monitoring of a product by a California manufacturer. **Choice of law may therefore require issue-specific analysis rather than the selection of one jurisdiction's law for the entire dispute merely because the final injury occurred there.**

This jurisdictional architecture is particularly important for AI products. Digital systems may be designed in one jurisdiction, commercially distributed from another, accessed globally, monitored through centralized corporate infrastructure, and allegedly cause physical or economic consequences thousands of kilometres away. Traditional territorial intuitions do not disappear, but they become substantially less capable of answering the legal question by themselves.

3.2. Eight Causes of Action: Responsibility Without Artificial Personhood

The British Columbia complaint pleads eight causes of action against all defendants: negligence for failure to warn law enforcement; negligent entrustment; aiding and abetting a mass shooting; negligence for failure to warn; negligent undertaking; negligence based upon design defect; strict product liability for design defect; and strict product liability for failure to warn.³³

Their doctrinal importance lies not merely in their number but in the object toward which they direct legal attribution.

³² Symeonides S. C. Choice of Law in Cross-Border Torts: Why Plaintiffs Win and Should // Hastings Law Journal. 2009. Vol. 61. P. 337–430.

³³ Cited Complaint.

The complaint does **not** ask the court to determine whether ChatGPT possesses intention, judgment, consciousness, autonomous legal will, or legal personality. Nor does it attempt to identify the artificial system itself as the bearer of tort responsibility. Instead, every cause of action returns attribution to identifiable human and corporate actors: the entities that allegedly designed the product, established its safety rules, monitored interactions, controlled access, made escalation decisions, issued or omitted warnings, and possessed the technical capacity to modify the system.

In that respect, the lawsuit is strikingly conventional.

The technology may be novel. The legal question is not.

It asks who possessed **knowledge, control, foreseeability, and preventive capacity**.

The first cause of action illustrates the point. The Crown and SD59 rely upon principles associated with *Tarasoff v. Regents of the University of California* and allege that the relationship between OpenAI and the user, combined with OpenAI's asserted knowledge of a specific danger, generated a duty to take reasonable protective measures, including notification of law enforcement. The complaint alleges that OpenAI's own systems identified the account, that human reviewers assessed the material, and that a referral to law enforcement was recommended but not implemented.

That theory is ambitious and its eventual acceptance cannot be assumed. *Tarasoff* arose from a different institutional relationship, and the extension of comparable principles to an AI provider presents an important doctrinal question. The relevance of the pleading lies precisely in the attempt to translate an established duty-to-protect concept into a new technological relationship without inventing an artificial legal subject.

The negligent-undertaking theory provides a different route. The plaintiffs allege that OpenAI voluntarily established safety functions for detecting dangerous interactions, escalating them to human reviewers, and potentially referring serious threats to authorities. Having allegedly undertaken those functions and publicly represented that such protective systems existed, the complaint argues that OpenAI incurred a duty to perform them with reasonable care.

This formulation is potentially more significant than a generalized proposition that every AI provider must continuously monitor every user. The complaint alleges something narrower and more concrete: **the manufacturer created a safety architecture, the architecture allegedly identified a particular risk, human reviewers allegedly recognized that risk, and the protective process was then allegedly interrupted by a corporate decision.**

The attribution chain therefore runs not from human actor to autonomous machine, but in the opposite direction—from the machine back through the safety system to the corporation that designed, supervised, and controlled it.

3.3. Tort Liability as a Form of AI Governance

The British Columbia litigation also belongs within a broader academic debate concerning whether traditional private law can function as a mechanism for governing emerging technologies.

Matthew U. Scherer identified a structural difficulty in AI regulation well before the present generation of generative systems became widely deployed. Rapid technological development, opacity, uncertain risk profiles, and the difficulty of specifying technical rules *ex ante* can make conventional prescriptive regulation comparatively slow and inflexible. Scherer therefore examined liability mechanisms as part of a wider regulatory strategy capable of influencing developer behaviour without requiring government to prescribe every engineering detail in advance.³⁴

That insight has direct relevance to *British Columbia v. Altman*.

The court is not being asked to write a neural-network architecture or determine the technical parameters of a language model. Rather, the plaintiffs seek to use ordinary doctrines of negligence and product liability to create incentives—and, through injunctive relief, potentially enforceable obligations—concerning monitoring, escalation, warnings, account controls, and safety procedures.

Private law, on this view, is not merely compensatory. It can also perform a **risk-allocation and behaviour-shaping function**.

Miriam C. Buiten develops the product-liability aspect of this problem more directly. Her analysis of defective AI proposes that allocation of responsibility should take account of the parties' relative **control over product risk and awareness of that risk**. AI systems complicate traditional liability relationships because users may possess limited knowledge of system architecture and limited capacity to control design-level risks, while producers may retain considerably greater knowledge and control over development, testing, updating, monitoring, and safeguards.³⁵

This framework maps closely onto the structure of the British Columbia complaint.

The allegations repeatedly concentrate upon matters allegedly remaining within the producer's sphere of control: model design, safety specifications, monitoring architecture, human-review procedures, account deactivation and re-access, warnings, escalation criteria, and managerial decisions concerning identified risks.

The doctrinal significance is considerable. AI complexity does not necessarily weaken the case for identifying a responsible actor. Under this approach, greater technical complexity may make the questions of **superior knowledge and superior control more important, not less important**.

³⁴ Scherer M. U. Regulating Artificial Intelligence Systems: Risks, Challenges, Competencies, and Strategies // Harvard Journal of Law & Technology. 2016. Vol. 29, No. 2. P. 353–400.

³⁵ Buiten M. C. Product Liability for Defective AI // European Journal of Law and Economics. 2024. Vol. 57. P. 239–273. DOI: 10.1007/s10657-024-09794-z.

That conclusion is also consistent with the producer-liability framework developed in *The Tool Has a Manufacturer: Producer Liability for Legal AI*. There, responsibility is treated as potentially cumulative rather than mutually exclusive: responsibility of a downstream professional or user does not automatically extinguish independent duties arising from the manufacturer's superior technical knowledge, architectural control, representations, and capacity to prevent systemic harm.³⁶

Although the factual setting of the present litigation is radically different from professional legal AI, the underlying attribution principle is closely related:

the absence of artificial voluntas does not create an absence of responsibility; it redirects the inquiry toward the human and corporate actors who designed, controlled, and deployed the instrument.

3.4. ChatGPT as a Product: The Design-Defect Theory

Perhaps the most consequential doctrinal move in the complaint is its explicit attempt to bring ChatGPT within **California strict products-liability law**.

The seventh cause of action alleges that ChatGPT constitutes a product and that the product was defectively designed. The complaint expressly seeks to locate the alleged defect in the system's design rather than in independent third-party content. It states, in unusually concise terms:

"ChatGPT performed as designed; the defect lay in that design."³⁷

The pleading further asserts that the claims do not rest upon treating OpenAI merely as the publisher of third-party speech. Instead, the plaintiffs characterize the relevant content as generated by the product itself and the alleged defect as residing in the design choices responsible for that behaviour.

That characterization is doctrinally important because it attempts to move the dispute away from a conventional intermediary-liability paradigm and toward **manufacturer responsibility for the operation of the product itself**.

The negligence-based design-defect count proceeds in parallel. The plaintiffs allege that feasible safer alternatives existed, including stronger refusal mechanisms, automated detection and escalation, mandatory human intervention in high-risk interactions, and responses designed to interrupt rather than sustain dangerous interactions. Again, these are pleaded allegations, not adjudicated facts. Their significance lies in the legal structure: the complaint identifies particular design choices and alleges the existence of alternative designs.

³⁶ Kildeev, A., *The Tool Has a Manufacturer: Producer Liability for Legal AI* (2026), SSRN Working Paper No. [7388178](https://ssrn.com/abstract=7388178), DOI: 10.2139/ssrn.7388178.

³⁷ Cited Complaint.

The eighth cause of action adds strict product liability for failure to warn. It alleges that the product was distributed without adequate warnings concerning risks that the plaintiffs contend were known or knowable to its manufacturer.

Whether a general-purpose generative AI system will ultimately be treated as a “product” for all relevant purposes of California strict-liability doctrine is therefore not a semantic question. It is potentially decisive. Classification affects the applicable liability architecture, the role of defect analysis, the relevance of reasonable alternative design, and the possibility of imposing duties independent of traditional negligence.

The larger theoretical question can be expressed simply:

If the system is treated as a product, who bears responsibility for defects embedded in its design?

The complaint’s answer is equally straightforward: not the probabilistic instrument itself, but those who allegedly designed, manufactured, distributed, monitored, and controlled it.

3.5. From Compensation to Preventive Architecture

The most important aspect of the British Columbia action may ultimately lie not in the damages claims but in the requested injunctive relief.

The Crown and SD59 seek compensatory damages for past and future losses, including alleged property damage and extraordinary public expenditures. They also request punitive damages, restitution, disgorgement, costs, interest, and declaratory relief.³⁸

But the prayer for relief goes substantially further.

The plaintiffs ask the federal court to require the defendants to implement operational safeguards addressing dangerous conversations; automated detection coupled with human-review escalation; referral procedures for credible threats; stronger warnings; effective account-level enforcement; auditable safety controls; and **quarterly compliance audits by an independent monitor**.

This is where the litigation changes institutional character.

Formally, the Province is acting as a civil plaintiff. It is **not exercising American regulatory authority**, and the lawsuit should not be confused with an administrative enforcement proceeding. Its claims remain tort and product-liability claims that defendants may contest on duty, defect, causation, standing, damages, jurisdictional, procedural, and other grounds.

Functionally, however, the requested injunction reaches beyond retrospective compensation. If such relief were granted, judicial enforcement would directly affect how a mass-market AI system is monitored and governed in the future.

³⁸ Cited Complaint.

The British Columbia government itself has framed the action in broader accountability terms. In announcing the lawsuit, Attorney General Niki Sharma expressly connected the proceeding with demands for stronger AI safeguards, greater transparency, and measures intended to prevent comparable harm.³⁹ The Province also drew an institutional analogy to its previous use of litigation against tobacco, opioid, and vaping companies, emphasizing the use of courts to prevent the costs of alleged corporate wrongdoing from being externalized onto the public.

That analogy should not be overstated. AI systems are not tobacco, opioids, or vaping products, and the substantive legal regimes are different. The importance of the comparison lies elsewhere: **British Columbia is explicitly presenting civil litigation as an instrument of corporate accountability and prospective risk control.**

That is the fundamental change represented by the third wave.

The first wave of litigation was primarily concerned with whether individual victims and families could recover for alleged injury. The second enlarged the group of private claimants. The Crown action introduces a public institutional plaintiff seeking not merely compensation for public losses but judicially enforceable changes in the manufacturer's future conduct.

The progression can therefore be stated more accurately as:

private injury → expanded private litigation → public-entity intervention → requested preventive controls over the manufacturer.

This is not yet regulation by judicial decree. No injunction has been granted, liability has not been established, and the defendants retain substantial procedural and substantive arguments.

But it is unmistakably **litigation aimed at governance.**

3.6. The Manufacturer Returns to the Centre of the Legal Inquiry

The broader significance of *British Columbia v. Altman* lies precisely in what the action does **not** require.

It does not require recognition of artificial consciousness.

It does not require electronic personhood.

It does not require the attribution of independent intention or judgment to an LLM.

It does not even require the court to resolve the philosophical question of whether the system “understands” what it generates.

³⁹ British Columbia. Ministry of Attorney General. Attorney General's Statement on B.C. Filing Legal Action Against OpenAI. Sept. 21, 2026. URL: <https://news.gov.bc.ca/releases/2026AG0067-001105> (accessed: 30.09.2026).

Traditional legal categories are sufficient to begin the inquiry: duty, knowledge, foreseeability, undertaking, product defect, warning, causation, control, and capacity to prevent harm.

The conceptual movement is therefore away from anthropomorphism and back toward attribution.

An AI system may produce the immediate output, but the product does not design its own commercial architecture, formulate corporate safety policies, determine escalation rules, establish account controls, issue public representations, decide what warnings consumers receive, or bear legal responsibility for institutional decisions. Those functions remain attributable to natural and juridical persons.

This is precisely where the Tumbler Ridge litigation intersects with the central proposition of *The Tool Has a Manufacturer*.

Technological complexity does not dissolve responsibility. It makes the **architecture of responsibility** more important.

The harder it becomes for an ordinary user to understand the internal operation of an AI system, the more consequential become questions of who possesses superior technical knowledge, who controls the relevant safeguards, who profits from deployment, who can alter the product, and who is institutionally capable of preventing recurrence.

The legal question is therefore not whether the artificial instrument should be blamed.

It is whether the existence of that instrument can become a mechanism through which the humans and corporations that manufacture and control it escape ordinary principles of attribution.

British Columbia v. Altman places that question squarely before an American federal court.

CHAPTER 4. CONCLUSION: FROM POST-MORTEM LITIGATION TO AN ARCHITECTURE OF RESPONSIBILITY

The significance of the litigation examined in this Article extends far beyond OpenAI, a single model, or even the extraordinary circumstances that gave rise to these proceedings. It concerns the emerging legal position of the entire class of companies that design, deploy, monitor, control and profit from artificial-intelligence systems.

The central legal question is not whether an AI system can itself bear responsibility. It cannot. The relevant questions remain fundamentally human and corporate. Who designed the system? Who determined its safety architecture? Who established the monitoring and escalation procedures? Who received information about an identified risk? Who possessed the authority to intervene? And, ultimately, who decided whether knowledge generated inside the system would be translated into action outside it?

These questions apply not merely to OpenAI. They apply, with factual variations, throughout Big Tech.

4.1. Responsibility Returns to the Manufacturer

The language surrounding artificial intelligence has often obscured responsibility by presenting technological systems as if they were independent actors. Models are said to “decide,” “understand,” “refuse,” “choose,” or “fail.” The more anthropomorphic the description becomes, the easier it is for the human and corporate actors behind the system to disappear from view.

Tort and product-liability law reverse that displacement.

Technical complexity does not create legal personality. Probabilistic output does not produce independent legal agency. Nor does the capacity of a system to generate language, detect patterns or initiate internal escalation transform that system into the bearer of legal responsibility.

The British Columbia Complaint instead relies upon familiar doctrines: negligence, negligent entrustment, aiding and abetting, negligent undertaking, design defect, failure to warn and strict product liability.

This is important. The litigation does not require the creation of an electronic person or an artificial bearer of duty. It asks whether existing doctrines of responsibility can be applied to actors possessing knowledge, control, authority and preventive capacity.

That inquiry returns inevitably to the manufacturer and controller.

A corporation that designs a system, establishes its safety mechanisms, controls its deployment, monitors its operation, receives internal warnings and retains authority over the response cannot make responsibility disappear merely by interposing an artificial system between itself and the consequences of its decisions.

The machine remains an instrument. The manufacturer remains a legal actor.

4.2. The Bitter Irony: The AI Did Not Fail

There is, however, a deeper and particularly bitter irony at the centre of this case.

On the allegations pleaded, the AI system did not fail at the critical detection stage.

The Complaint alleges that automated monitoring identified the relevant account and escalated it for human review. Nor, according to the pleading, did the human safety personnel fail at the next stage. Human reviewers allegedly assessed the risk as credible and recommended referral to the Royal Canadian Mounted Police. The Complaint further alleges that the recommendation was not acted upon at the corporate decision-making level.

If those allegations are substantiated, this is not principally a story of artificial intelligence escaping human control. It is almost the opposite. The instrument allegedly detected. Human reviewers allegedly assessed. The internal safety architecture allegedly escalated. The asserted breakdown occurred only after the information reached the level at which corporate authority had to be exercised. This conclusion concerns the alleged

breakdown in the internal detection-and-escalation chain; it does not itself establish legal causation or resolve the effects of intervening acts.

The significance of that distinction is difficult to overstate.

Much of the contemporary debate about AI safety is framed around the possibility that increasingly sophisticated systems may become uncontrollable. This litigation presents a different and potentially more immediate problem: **the failure to exercise control by the very actor in whom control already resides.**

In the terminology of the Architecture of Responsibility, this may be understood as *abdicatio officii domini*: the failure of the controller to perform the non-delegable function attached to its own authority.

The relevant failure, if proved, was therefore not artificial autonomy. It was the alleged non-exercise of human corporate authority.

Control cannot mean merely possessing the technical ability to monitor, suspend, modify or terminate an AI system. It must also include the responsibility to act when the system successfully delivers information requiring human judgment.

A corporation cannot claim the benefits of control when deployment, commercialization and profit are concerned, while treating itself as a passive observer when its own safety architecture identifies a serious risk.

Authority, control and responsibility must remain connected.

The machine remained an instrument. The human reviewers remained capable of judgment. The alleged failure occurred where ultimate authority resided.

4.3. The Consequences for Big Tech

This distinction carries consequences far beyond the defendants named in the present litigation.

The first consequence concerns evidence.

Internal AI-safety architecture is likely to become increasingly important in litigation involving consequential AI systems. Detection mechanisms, safety thresholds, escalation procedures, model logs, human-review records, internal communications, risk assessments and management decisions may become central evidence in determining what a manufacturer knew, when it knew it, what actions were technically available and why those actions were or were not taken. AI liability therefore cannot be reduced to the isolated text produced by a chatbot. The legally relevant object is the larger socio-technical system: the model, the monitoring architecture, the human reviewers, the escalation structure, the decision-makers and the corporate policies connecting them.

The second consequence concerns corporate organization.

Large technology companies frequently operate through complex structures involving parent corporations, subsidiaries, foundations, operating entities, contractors and distributed technical teams. Such structures may be commercially legitimate, but they should not permit responsibility to evaporate between entities whenever a serious safety event occurs. Where knowledge exists somewhere within the corporate system and authority exists elsewhere, the legal question will increasingly become whether the organization was reasonably designed to ensure that knowledge reached the person or body capable of acting upon it.

The third consequence concerns the relationship between tort liability and product architecture.

The British Columbia Complaint does not seek damages alone. It also asks for prospective safeguards relating to violent-content detection, human escalation, referral to law enforcement, account enforcement, warnings, auditable controls and independent compliance monitoring. That represents an important shift. AI litigation may increasingly concern not simply compensation for past harm, but judicially compelled alteration of the architecture through which future risks are detected and managed.

For Big Tech, that is qualitatively different from ordinary damages litigation.

The court may ultimately determine that some or all of these theories cannot be sustained. That remains a matter for adjudication. But the questions themselves will not disappear. Similar questions will arise wherever companies deploy systems capable of generating consequential interactions while simultaneously maintaining monitoring, moderation and intervention mechanisms.

The proposition developed in *The Tool Has a Manufacturer* therefore acquires a broader significance.

The manufacturer is not responsible because the machine possesses a will.

The manufacturer may be responsible precisely because the machine does not.

Responsibility follows the actors who design, control, monitor, authorize, profit and possess the capacity to prevent foreseeable harm.

4.4. Courts Are the Backstop, Not the Safety System

There is nevertheless a fundamental weakness in relying upon litigation as the principal mechanism of AI governance.

Courts generally encounter the safety architecture after something has already gone wrong. They reconstruct events. They examine evidence. They determine duty, breach, causation and damages. They may impose liability and order prospective relief. These functions are indispensable. But they are inherently retrospective. In that sense, litigation is a form of legal post-mortem.

In the present litigation, that description carries an especially painful literal dimension. Whatever the eventual disposition of the claims, no judgment can reverse the events that gave rise to them. That is precisely why AI safety cannot be left primarily to tort litigation.

The purpose of the courts is to determine responsibility after a legally cognizable controversy arises. The purpose of regulation should be to reduce the probability that the controversy arises at all.

The litigation does not itself compel any particular regulatory model; it does, however, expose recurring governance questions that tort adjudication alone can address only retrospectively

The United States therefore requires a coherent federal baseline for consequential AI systems. Such regulation need not prescribe every technical parameter of every model. Nor should it attempt to freeze technological development through rigid engineering mandates. Its central function should instead be to establish a minimum architecture of responsibility across the industry.

That architecture should require documented procedures for detection and escalation of serious risks; preservation of safety-relevant logs and internal decision records; clear assignment of human responsibility for high-risk decisions; age-sensitive safeguards where systems interact with minors; defined procedures for responding to credible and imminent threats; independent auditing of critical safety controls; serious-incident reporting; and corporate structures that permit the identification of the person or body possessing final authority.

It should also establish an essential principle: compliance with regulatory minimums should not automatically extinguish ordinary tort and product-liability responsibility where unreasonable conduct can otherwise be proved.

The objective is not to create a federal morality code for artificial intelligence.

It is to answer much more concrete questions.

1. Who must monitor?
2. Who must review?
3. Who must decide?
4. What must be recorded?
5. When must an incident be escalated?
6. Who has authority to act?
7. And who remains accountable when the system fails?

Without a federal baseline, these questions will continue to be answered piecemeal through individual lawsuits, inconsistent corporate practices and judicial reconstruction after harm has already occurred. That is neither an efficient regulatory structure nor a sufficient safety architecture.

Tort litigation must therefore remain available, but as the **backstop of responsibility**, not as the principal mechanism through which society discovers what AI manufacturers should have done.

The distinction is ultimately one between ex post accountability and ex ante prevention.

The first belongs primarily to courts.

The second requires regulation, engineering and an architecture of responsibility before the harm occurs.

The litigation discussed in this Article may succeed, fail, narrow, settle, or ultimately be resolved on grounds substantially different from the broader issues considered here. Those questions belong to the court.

But the structural lesson already extends beyond this particular dispute.

Artificial intelligence does not eliminate human responsibility. Technological complexity does not dissolve corporate responsibility. And no degree of sophistication in an artificial system transforms the instrument into the dominus.

The deepest lesson of this case may therefore be the most uncomfortable one.

The principal danger is not always the loss of control over AI. Sometimes it is the abdication of control by those who already possess it.

EXHIBIT A

LITIGATION MAP: OPENAI / CHATGPT HARM CLAIMS

Florida State University and Tumbler Ridge litigation — docket-focused factual appendix

Scope and status. This exhibit is designed as a factual docket map, not as a substitute for the article's legal analysis. All statements below are allegations unless expressly identified as docket facts; no allegation is presented as a judicial finding. Docket information is current through 29 September 2026. The Northern District of California Tumbler Ridge actions are presented in waves to make the litigation sequence visible without reproducing the narrative discussion in the article.

A. Florida State University litigation

Case / filing	Plaintiff(s)	Defendant(s)	Allegations / causes pleaded (summary)
Joshi v. OpenAI Foundation et al. No. 4:26-cv-00222 N.D. Fla. Filed 10 May 2026	Vandana Joshi	OpenAI Foundation and affiliated OpenAI entities; Phoenix Ikner	• Product-liability and negligence theories arising from ChatGPT's alleged role in the events preceding the FSU shooting. • Allegation that ChatGPT materially assisted the attacker's planning and decision-making. • Design-defect / failure-to-warn theories directed at the product and its safeguards. • Wrongful-death and related damages sought by the victim's family.

B. Tumbler Ridge — Wave I: seven family actions filed 29 April 2026

Common core alleged across the April family complaints: OpenAI allegedly detected dangerous activity months before the attack; internal safety personnel allegedly recommended a law-enforcement referral; leadership allegedly did not notify the RCMP; the user allegedly regained access after account deactivation; plaintiffs plead negligence, negligent entrustment/undertaking, failure-to-warn and design-defect/product-liability theories, together with aiding-and-abetting and individualized wrongful-death, survival, or personal-injury claims.

Case / filing	Plaintiff(s)	Defendant(s)	Allegations / causes pleaded (summary)
Stacey v. Altman et al. No. 3:26-cv-03701-JSC N.D. Cal.	Mark Stacey	Samuel Altman; OpenAI Foundation; OpenAI OpCo, LLC; OpenAI Group PBC	Wave I common core; wrongful-death/survival claims tied to the death of a school employee.
Schofield et al. v. Altman et al. No. 3:26-cv-03702-JSC N.D. Cal.	Timm Schofield; Rebekah Schofield	Samuel Altman; OpenAI Foundation; OpenAI OpCo, LLC; OpenAI Group PBC	Wave I common core; wrongful-death/survival claims on behalf of their child.
Mwansa, Sr. et al. v. Altman et al. No. 3:26-cv-03703-JSC N.D. Cal.	Abel Mwansa, Sr.; Bwalya Chisanga	Samuel Altman; OpenAI Foundation; OpenAI OpCo, LLC; OpenAI Group PBC	Wave I common core; wrongful-death/survival and product-liability claims on behalf of their child.
M.G. et al. v. Altman et al. No. 3:26-cv-03704-JSC N.D. Cal.	M.G.; Cia Edmonds	Samuel Altman; OpenAI Foundation; OpenAI OpCo, LLC; OpenAI Group PBC	Wave I common core; personal-injury / product-liability claims associated with the Tumbler Ridge attack.
Hayer v. Altman et al. No. 3:26-cv-03705-JSC N.D. Cal.	Lori Hayer	Samuel Altman; OpenAI Foundation; OpenAI OpCo, LLC; OpenAI Group PBC; docket also lists Quintin Benoit	Wave I common core; wrongful-death/survival claims on behalf of her child.

Case / filing	Plaintiff(s)	Defendant(s)	Allegations / causes pleaded (summary)
Lampert v. Altman et al. No. 3:26-cv-03706-JSC N.D. Cal.	Sarah Lampert	Samuel Altman; OpenAI Foundation; OpenAI OpCo, LLC; OpenAI Group PBC	Wave I common core; wrongful-death/survival claims on behalf of her child.
Younge et al. v. Altman et al. No. 3:26-cv-03707-JSC N.D. Cal.	Lance Younge; Jennifer Geary	Samuel Altman; OpenAI Foundation; OpenAI OpCo, LLC; OpenAI Group PBC	Wave I common core; personal-injury / product-liability claims associated with the Tumbler Ridge attack.

C. Tumbler Ridge — Wave II: thirty survivor / educator actions filed beginning 2 September 2026

Common allegation profile for the thirty Wave II actions: (1) failure to warn law enforcement after an alleged internal safety escalation; (2) negligent entrustment / negligent undertaking; (3) alleged design defect and inadequate warnings; (4) strict-product-liability theories; (5) aiding-and-abetting theory; and (6) compensatory and punitive damages based on claimed physical and/or emotional injury. The cases below are listed individually because each has a separate federal docket, even though the factual core substantially overlaps.

Case / filing	Plaintiff(s)	Defendant(s)	Allegations / causes pleaded (summary)
A.C. v. Altman et al. No. 3:26-cv-09342-JSC N.D. Cal.	A.C.	Samuel Altman; OpenAI Foundation; OpenAI OpCo, LLC; OpenAI Group PBC	Wave II common allegation profile (see paragraph above).
D.I. v. Altman et al. No. 3:26-cv-09343-JSC N.D. Cal.	D.I.	Samuel Altman; OpenAI Foundation; OpenAI OpCo, LLC; OpenAI Group PBC	Wave II common allegation profile (see paragraph above).
Deidre Rushlow v. Altman et al. No. 3:26-cv-09344-JSC N.D. Cal.	Deidre Rushlow	Samuel Altman; OpenAI Foundation; OpenAI OpCo, LLC; OpenAI Group PBC	Wave II common allegation profile (see paragraph above).
Hodgkinson v. Altman et al. No. 3:26-cv-09345-JSC N.D. Cal.	Hodgkinson	Samuel Altman; OpenAI Foundation; OpenAI OpCo, LLC; OpenAI Group PBC	Wave II common allegation profile (see paragraph above).
Bertrand v. Altman et al. No. 3:26-cv-09346-JSC N.D. Cal.	Bertrand	Samuel Altman; OpenAI Foundation; OpenAI OpCo, LLC; OpenAI Group PBC	Wave II common allegation profile (see paragraph above).
Geer v. Altman et al. No. 3:26-cv-09347-JSC N.D. Cal.	Geer	Samuel Altman; OpenAI Foundation; OpenAI OpCo, LLC; OpenAI Group PBC	Wave II common allegation profile (see paragraph above).
Maitland v. Altman et al. No. 3:26-cv-09348-JSC N.D. Cal.	Maitland	Samuel Altman; OpenAI Foundation; OpenAI OpCo, LLC; OpenAI Group PBC	Wave II common allegation profile (see paragraph above).
Tardif v. Altman et al. No. 3:26-cv-09349-JSC N.D. Cal.	Tardif	Samuel Altman; OpenAI Foundation; OpenAI OpCo, LLC; OpenAI Group PBC	Wave II common allegation profile (see paragraph above).

Case / filing	Plaintiff(s)	Defendant(s)	Allegations / causes pleaded (summary)
Zatwarniski v. Altman et al. No. 3:26-cv-09350-JSC N.D. Cal.	Zatwarniski	Samuel Altman; OpenAI Foundation; OpenAI OpCo, LLC; OpenAI Group PBC	Wave II common allegation profile (see paragraph above).
Pilgrim v. Altman et al. No. 3:26-cv-09351-JSC N.D. Cal.	Pilgrim	Samuel Altman; OpenAI Foundation; OpenAI OpCo, LLC; OpenAI Group PBC	Wave II common allegation profile (see paragraph above).
Pigeon v. Altman et al. No. 3:26-cv-09352-JSC N.D. Cal.	Pigeon	Samuel Altman; OpenAI Foundation; OpenAI OpCo, LLC; OpenAI Group PBC	Wave II common allegation profile (see paragraph above).
Lenart v. Altman et al. No. 3:26-cv-09353-JSC N.D. Cal.	Lenart	Samuel Altman; OpenAI Foundation; OpenAI OpCo, LLC; OpenAI Group PBC	Wave II common allegation profile (see paragraph above).
Cameron v. Altman et al. No. 3:26-cv-09354-JSC N.D. Cal.	Cameron	Samuel Altman; OpenAI Foundation; OpenAI OpCo, LLC; OpenAI Group PBC	Wave II common allegation profile (see paragraph above).
Norbury v. Altman et al. No. 3:26-cv-09355-JSC N.D. Cal.	Norbury	Samuel Altman; OpenAI Foundation; OpenAI OpCo, LLC; OpenAI Group PBC	Wave II common allegation profile (see paragraph above).
Gruntman v. Altman et al. No. 3:26-cv-09356-JSC N.D. Cal.	Gruntman	Samuel Altman; OpenAI Foundation; OpenAI OpCo, LLC; OpenAI Group PBC	Wave II common allegation profile (see paragraph above).
Pisarski v. Altman et al. No. 3:26-cv-09357-JSC N.D. Cal.	Pisarski	Samuel Altman; OpenAI Foundation; OpenAI OpCo, LLC; OpenAI Group PBC	Wave II common allegation profile (see paragraph above).
E.R. v. Altman et al. No. 3:26-cv-09358-JSC N.D. Cal.	E.R.	Samuel Altman; OpenAI Foundation; OpenAI OpCo, LLC; OpenAI Group PBC	Wave II common allegation profile (see paragraph above).
G.L. v. Altman et al. No. 3:26-cv-09359-JSC N.D. Cal.	G.L.	Samuel Altman; OpenAI Foundation; OpenAI OpCo, LLC; OpenAI Group PBC	Wave II common allegation profile (see paragraph above).
A.R. v. Altman et al. No. 3:26-cv-09360-JSC N.D. Cal.	A.R.	Samuel Altman; OpenAI Foundation; OpenAI OpCo, LLC; OpenAI Group PBC	Wave II common allegation profile (see paragraph above).
C.W. v. Altman et al. No. 3:26-cv-09362-	C.W.	Samuel Altman; OpenAI Foundation; OpenAI OpCo, LLC; OpenAI Group PBC	Wave II common allegation profile (see paragraph above).

Case / filing	Plaintiff(s)	Defendant(s)	Allegations / causes pleaded (summary)
JSC N.D. Cal.			
A.N. et al. v. Altman et al. No. 3:26-cv-09363-JSC N.D. Cal.	A.N. et al.	Samuel Altman; OpenAI Foundation; OpenAI OpCo, LLC; OpenAI Group PBC	Wave II common allegation profile (see paragraph above).
R.R. v. Altman et al. No. 3:26-cv-09366-JSC N.D. Cal.	R.R.	Samuel Altman; OpenAI Foundation; OpenAI OpCo, LLC; OpenAI Group PBC	Wave II common allegation profile (see paragraph above).
G.M. v. Altman et al. No. 3:26-cv-09367-JSC N.D. Cal.	G.M.	Samuel Altman; OpenAI Foundation; OpenAI OpCo, LLC; OpenAI Group PBC	Wave II common allegation profile (see paragraph above).
L.N. v. Altman et al. No. 3:26-cv-09368-JSC N.D. Cal.	L.N.	Samuel Altman; OpenAI Foundation; OpenAI OpCo, LLC; OpenAI Group PBC	Wave II common allegation profile (see paragraph above).
Poirier v. Altman et al. No. 3:26-cv-09369-JSC N.D. Cal.	Poirier	Samuel Altman; OpenAI Foundation; OpenAI OpCo, LLC; OpenAI Group PBC	Wave II common allegation profile (see paragraph above).
M.B. et al. v. Altman et al. No. 3:26-cv-09370-JSC N.D. Cal.	M.B. et al.	Samuel Altman; OpenAI Foundation; OpenAI OpCo, LLC; OpenAI Group PBC	Wave II common allegation profile (see paragraph above).
O.H. et al. v. Altman et al. No. 3:26-cv-09372-JSC N.D. Cal.	O.H. et al.	Samuel Altman; OpenAI Foundation; OpenAI OpCo, LLC; OpenAI Group PBC	Wave II common allegation profile (see paragraph above).
B.W. et al. v. Altman et al. No. 3:26-cv-09373-JSC N.D. Cal.	B.W. et al.	Samuel Altman; OpenAI Foundation; OpenAI OpCo, LLC; OpenAI Group PBC	Wave II common allegation profile (see paragraph above).
Hoekstra et al. v. Altman et al. No. 3:26-cv-09377-JSC N.D. Cal.	Hoekstra et al.	Samuel Altman; OpenAI Foundation; OpenAI OpCo, LLC; OpenAI Group PBC	Wave II common allegation profile (see paragraph above).
Deeley et al. v. Altman et al. No. 3:26-cv-09378-JSC N.D. Cal.	Deeley et al.	Samuel Altman; OpenAI Foundation; OpenAI OpCo, LLC; OpenAI Group PBC	Wave II common allegation profile (see paragraph above).

D. Public-entity action — Province of British Columbia and School District No. 59

Case / filing	Plaintiff(s)	Defendant(s)	Allegations / causes pleaded (summary)
His Majesty the King in Right of the Province of British Columbia et al. v. Altman et al. No. 4:26-cv-10743-JSC N.D. Cal. Filed 21 Sept. 2026	His Majesty the King in Right of the Province of British Columbia; Board of Education of School District No. 59 (Peace River South)	Samuel Altman; OpenAI Foundation (f/k/a OpenAI, Inc.); OpenAI OpCo, LLC; OpenAI Holdings, LLC; OpenAI Group PBC	• Negligence — alleged failure to warn law enforcement. • Negligent entrustment. • Aiding and abetting a mass shooting. • Negligence — failure to warn. • Negligent undertaking. • Negligence — design defect. • Strict product liability — design defect. • Strict product liability — failure to warn. • Plaintiffs seek compensatory and punitive damages and injunctive relief; the Province and SD59 also allege extraordinary public-response and property-related losses.

Procedural note

The seven April Tumbler Ridge actions were treated as related in June 2026, with Stacey v. Altman serving as the lowest-numbered lead action. By late September 2026, the Northern District of California docket reflected the thirty Wave II actions and the Province/SD59 action before Judge Jacqueline Scott Corley. The Province's filed complaint identifies Case No. 3:26-cv-10743 on its caption; an earlier ECF header displayed a different divisional/judicial suffix before subsequent reassignment. For this exhibit, the current docket form 3:26-cv-10743-JSC is used.

Source note

Primary / docket sources used for compilation: filed complaint in His Majesty the King in Right of the Province of British Columbia and Board of Education of School District No. 59 (Peace River South) v. Altman et al.; U.S. District Court for the Northern District of California public docket/new-filings records for the related Tumbler Ridge actions; U.S. District Court for the Northern District of Florida docket in Joshi v. OpenAI Foundation et al., No. 4:26-cv-00222; and publicly available federal docket summaries for the April family actions. The exhibit deliberately summarizes allegations and omits extended quotations from pleadings.

References

1. Kildeev, A., Disney and Warner Brothers Discovery Lawsuits against Midjourney and Matthew and Maria Raine v. OpenAI and Sam Altman and their Implications for the Development of the Artificial Intelligence Industry. *Scientific Platform XXI Century*, Issue 4 (July 2025), pp. 20-48; DOI: 10.2139/ssrn.5856085, accessed on 28.09.26.
2. https://www.reuters.com/sustainability/boards-policy-regulation/openai-altman-sued-over-chatgpts-role-california-teens-suicide-2025-08-26/?utm_source=chatgpt.com;
3. <https://www.cnn.com/2025/08/26/tech/openai-chatgpt-teen-suicide-lawsuit>;
4. https://www.theguardian.com/technology/2025/aug/27/chatgpt-scrutiny-family-teen-killed-himself-sue-open-ai?utm_source=chatgpt.com;
5. https://www.techpolicy.press/breaking-down-the-lawsuit-against-openai-over-teens-suicide/?utm_source=chatgpt.com.
6. <https://dockets.justia.com/docket/california/candce/4:2026cv10743/479418>
7. https://www.techpolicy.press/breaking-down-the-lawsuit-against-openai-over-teens-suicide/?utm_source=chatgpt.com
8. Kildeev, A., *The Tool Has a Manufacturer: Producer Liability for Legal AI* (2026), SSRN Working Paper No. 7388178, DOI: 10.2139/ssrn.7388178.
9. Royal Canadian Mounted Police. Update: RCMP publicly confirm victims of Tumbler Ridge shooting. Tumbler Ridge, British Columbia, 12 February 2026. RCMP source.
10. Royal Canadian Mounted Police. Update: Tumbler Ridge fatal shootings. Tumbler Ridge, British Columbia, 11 February 2026. RCMP source.
11. Judd A. Tumbler Ridge shooter's ChatGPT activity flagged internally 7 months before tragedy // Global News. 20 February 2026; updated 21 February 2026. Global News article.
12. Boynton S. OpenAI says Tumbler Ridge shooter would be flagged to police today // Global News. 26 February 2026. Global News article.
13. Gebala et al. v. OpenAI Foundation et al., Notice of Civil Claim, Supreme Court of British Columbia, Vancouver Registry, No. S-261734, filed 9 March 2026. Notice of Civil Claim.
14. Urquhart C., Judd A. Family of Tumbler Ridge shooting survivor sues OpenAI // Global News. 9 March 2026. Global News article.
15. Stacey v. Altman et al., Case No. 3:26-cv-03701 // United States District Court for the Northern District of California. Filed Apr. 29, 2026. URL: <https://dockets.justia.com/docket/california/candce/3%3A2026cv03701/468851> (accessed: 29.09.2026).
16. Schofield et al. v. Altman et al., Case No. 3:26-cv-03702 // United States District Court for the Northern District of California. Filed Apr. 29, 2026. URL: <https://dockets.justia.com/docket/california/candce/3%3A2026cv03702/468852> (accessed: 29.09.2026).

17. Mwansa, Sr. et al. v. Altman et al., Case No. 3:26-cv-03703 // United States District Court for the Northern District of California. Filed Apr. 29, 2026. URL: <https://dockets.justia.com/docket/california/candce/3%3A2026cv03703/468853> (accessed: 29.09.2026).
18. M.G. et al. v. Altman et al., Case No. 3:26-cv-03704 // United States District Court for the Northern District of California. Filed Apr. 29, 2026. URL: <https://dockets.justia.com/docket/california/candce/3%3A2026cv03704/468854> (accessed: 29.09.2026).
19. Hayer v. Altman et al., Case No. 3:26-cv-03705 // United States District Court for the Northern District of California. Filed Apr. 29, 2026. URL: <https://dockets.justia.com/docket/california/candce/3%3A2026cv03705/468855> (accessed: 29.09.2026).
20. Lampert v. Altman et al., Case No. 3:26-cv-03706 // United States District Court for the Northern District of California. Filed Apr. 29, 2026. URL: <https://dockets.justia.com/docket/california/candce/3%3A2026cv03706/468856> (accessed: 29.09.2026).
21. Younge et al. v. Altman et al., Case No. 3:26-cv-03707 // United States District Court for the Northern District of California. Filed Apr. 29, 2026. URL: <https://dockets.justia.com/docket/california/candce/3%3A2026cv03707/468857> (accessed: 29.09.2026).
22. British Columbia. Ministry of Attorney General. Attorney General’s Statement on New Legal Filings Related to Tumbler Ridge Tragedy. Sept. 2, 2026. URL: <https://news.gov.bc.ca/releases/2026AG0061-001022> (accessed: 29.09.2026).
23. British Columbia. Ministry of Attorney General. Attorney General’s Statement on B.C. Filing Legal Action Against OpenAI. Sept. 21, 2026. URL: <https://news.gov.bc.ca/releases/2026AG0067-001105> (accessed: 29.09.2026).
24. His Majesty the King in Right of the Province of British Columbia and Board of Education of School District No. 59 (Peace River South) v. Samuel Altman et al. Complaint // U.S. District Court for the Northern District of California. Filed Sept. 21, 2026. 39 p. PACER, Doc 1.
25. Piper Aircraft Co. v. Reyno, 454 U.S. 235 (1981) // Supreme Court of the United States. URL: <https://www.law.cornell.edu/supremecourt/text/454/235> (accessed: 30.09.2026).
26. Sinochem International Co. Ltd. v. Malaysia International Shipping Corp., 549 U.S. 422 (2007) // Supreme Court of the United States. URL: <https://www.law.cornell.edu/supremecourt/text/06-102> (accessed: 30.09.2026).
27. Klaxon Co. v. Stentor Electric Manufacturing Co., 313 U.S. 487 (1941) // Supreme Court of the United States. URL: <https://www.law.cornell.edu/supremecourt/text/313/487> (accessed: 30.09.2026).
28. Hurtado v. Superior Court, 11 Cal. 3d 574, 522 P.2d 666 (1974) // Supreme Court of California. URL: <https://law.justia.com/cases/california/supreme-court/3d/11/574.html> (accessed: 30.09.2026).
29. Symeonides S. C. Choice of Law in Cross-Border Torts: Why Plaintiffs Win and Should // Hastings Law Journal. 2009. Vol. 61. P. 337–430.

30. Scherer M. U. Regulating Artificial Intelligence Systems: Risks, Challenges, Competencies, and Strategies // Harvard Journal of Law & Technology. 2016. Vol. 29, No. 2. P. 353–400.
31. Buiten M. C. Product Liability for Defective AI // European Journal of Law and Economics. 2024. Vol. 57. P. 239–273. DOI: 10.1007/s10657-024-09794-z.
32. British Columbia. Ministry of Attorney General. Attorney General’s Statement on B.C. Filing Legal Action Against OpenAI. Sept. 21, 2026. URL: <https://news.gov.bc.ca/releases/2026AG0067-001105> (accessed: 30.09.2026).